

Kabir_Notebook_6

November 27, 2024

1 Intro ML Homework 6

1.1 Name: Jaskin Kabir

1.2 Student ID: 801186717

Github: https://github.com/jaskinkabir/Intro_ML/tree/main/HM5

```
[63]: import pandas as pd
import numpy as np
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
import torch
from torch import nn

path = 'housing.csv'
housing = pd.DataFrame(pd.read_csv(path))

varlist = ['mainroad', 'guestroom', 'basement', 'hotwaterheating', 'airconditioning', 'prefarea', 'furnishingstatus']

# Defining the map function
def binary_map(x):
    return x.map({'yes': 1, 'no': 0, 'unfurnished': 0, 'semi-furnished': 1, 'furnished': 2})

# Applying the function to the housing list
housing[varlist] = housing[varlist].apply(binary_map)
housing.head()

df_train, df_test = train_test_split(housing, train_size=0.8, test_size=0.2, random_state=100)

Y_train_tensor_p1 = df_train.pop('price')
X_train = df_train

Y_test_tensor_p1 = df_test.pop('price')
X_test_batch = df_test
```

```

X_train.head()

b_vars = ['area', 'bedrooms', 'bathrooms', 'stories', 'mainroad', 'guestroom', 'basement', 'hotwaterheating', 'airconditioning', 'parking', 'prefarea']
X_p3 = housing[b_vars]
X_train_df_p3 = X_train[b_vars]
X_test_df_p3 = X_test_batch[b_vars]

X_train_tensor_p1 = X_train_df_p3.to_numpy()
X_test_tensor_p1 = X_test_df_p3.to_numpy()

scaler = preprocessing.StandardScaler().fit(X_train_tensor_p1)
X_train_tensor_p1 = scaler.transform(X_train_tensor_p1)
X_test_tensor_p1 = scaler.transform(X_test_tensor_p1)

```

```

[64]: from sklearn.metrics import root_mean_squared_error
import matplotlib.pyplot as plt

class Regressor(nn.Module):
    @classmethod
    def compare_results(cls, results1, results2):
        print(100 * (results1 - results2) / results1)

    def __init__(self, in_dim, out_dim, hidden_layers=[64,32], activation=nn.Tanh,):
        super().__init__()
        self.hidden_layers = hidden_layers
        self.activation = activation
        self.input_cols = []
        self.output_cols = []

        #Error Mode is a 5 bit integer, with each bit representing a feature
        # If the bit is 1, the feature is errored
        output_features = out_dim
        input_features = in_dim

        self.stack_list = [nn.Linear(input_features, hidden_layers[0]),
        activation()]
        for i in range(1, len(hidden_layers)):
            self.stack_list.extend([nn.Linear(hidden_layers[i-1],
        hidden_layers[i]), activation()])
        self.stack_list.extend([nn.Linear(hidden_layers[-1], output_features)])
        self.stack = nn.Sequential(*self.stack_list)

```

```

def train(self, epochs, X_train, X_test, Y_train, Y_test, alpha=1e-2,
↪ loss_fn=nn.MSELoss(),):

    val_hist = np.zeros(epochs)
    train_hist = np.zeros(epochs)

    optimizer = torch.optim.Adam(self.parameters(), lr=alpha)
    for i in range(epochs):
        optimizer.zero_grad()
        Y_pred = self.forward(X_train)
        loss = loss_fn(Y_pred.squeeze(), Y_train)
        loss.backward()
        optimizer.step()
        train_hist[i] = np.sqrt(loss.item())

        with torch.no_grad():
            Y_pred_val = self.forward(X_test)
            val_hist[i] = np.sqrt(loss_fn(Y_pred_val, Y_test).item())

    self.last_test = Y_test_tensor_p1
    self.last_pred = self.forward(X_test)
    self.last_score = val_hist[-1]

    self.last_epochs = epochs
    self.last_val_hist = val_hist
    self.last_train_hist = train_hist

def plot_loss(self, title):

    plt.plot(range(self.last_epochs), self.last_val_hist, label='Validation_
↪ Loss')
    plt.plot(range(self.last_epochs), self.last_train_hist, label='Training_
↪ Loss')
    plt.title(title)
    plt.xlabel('Epoch')
    plt.ylabel('RMS Loss')
    plt.legend()
    plt.show()

def forward(self, x):
    return self.stack(x)

def print_results(self):
    if self.last_score is None:

```

```

        raise ValueError('No results to print')
    print(f'MSE: {self.last_score:.2E}')

class CustomMSELoss(nn.Module):
    def __init__(self, lambda_val=0.0):
        super(CustomMSELoss, self).__init__()
        self.lambda_val = lambda_val

    def forward(self, predictions, targets):

        m = targets.size(0)

        errors = predictions - targets
        mse_loss = (1 / (2*m)) * torch.sum(errors ** 2)

        total_loss = mse_loss

    return total_loss

```

```

[65]: device = 'cpu'

X_train_tensor_p1 = torch.tensor(X_train_tensor_p1).to(device).float()
X_test_tensor_p1 = torch.tensor(X_test_tensor_p1).to(device).float()

Y_train_tensor_p1 = torch.tensor(Y_train_tensor_p1.to_numpy()).to(device).
    ↪float().view(-1, 1)
Y_test_tensor_p1 = torch.tensor(Y_test_tensor_p1.to_numpy()).to(device).float().
    ↪view(-1, 1)

model_1a = Regressor(
    in_dim=X_train_tensor_p1.shape[1],
    out_dim=1,
    hidden_layers=[8],
    activation=nn.ReLU,
).to(device)

model_1a.train(
    epochs=5000,
    X_train=X_train_tensor_p1,
    X_test=X_test_tensor_p1,
    Y_train=Y_train_tensor_p1,
    Y_test=Y_test_tensor_p1,

```

```

    alpha=1e-1,
    loss_fn=nn.MSELoss(),
)

model_1a.plot_loss('Problem 1a: 1 Hidden Layer Loss')
model_1a.print_results()


```

/home/jaskin/.local/lib/python3.11/site-packages/torch/nn/modules/loss.py:608: UserWarning: Using a target size (torch.Size([436, 1])) that is different to the input size (torch.Size([436])). This will likely lead to incorrect results due to broadcasting. Please ensure they have the same size.

```

    return F.mse_loss(input, target, reduction=self.reduction)

```



MSE: 2.09E+06

```

[67]: from sklearn.datasets import load_breast_cancer

breast = load_breast_cancer()
X_2 = breast.data
Y_2 = breast.target

```

```

X_train_2, X_test_2, Y_train_2, Y_test_2 = train_test_split(X_2, Y_2,
↳train_size=0.8, test_size=0.2, random_state=00)

scaler = preprocessing.StandardScaler().fit(X_train_2)
X_train_2 = scaler.transform(X_train_2)
X_test_2 = scaler.transform(X_test_2)

```

```

[68]: from sklearn.metrics import accuracy_score, precision_score, recall_score,
↳f1_score, confusion_matrix, classification_report

class Classifier(nn.Module):
    @classmethod
    def compare_results(cls, results1, results2):
        print('Comparing results:')
        comparisons = {
            'accuracy': 100*(results1['accuracy'] - results2['accuracy'])/
↳results1['accuracy'],
            'precision': 100*(results1['precision'] - results2['precision'])/
↳results1['precision'],
            'recall': 100*(results1['recall'] - results2['recall'])/
↳results1['recall'],
            'f1': 100*(results1['f1'] - results2['f1'])/results1['f1']
        }
        for key, value in comparisons.items():
            print(f'{key}: {value} %')

    def __init__(self):
        super().__init__()

    def get_results(self, Y_test=None, Y_pred=None):
        if Y_test is None:
            Y_test = self.last_test
        if Y_pred is None:
            Y_pred = self.last_pred

        if isinstance(Y_test, torch.Tensor):
            Y_test = Y_test.cpu().detach().numpy()
        if isinstance(Y_pred, torch.Tensor):
            Y_pred = Y_pred.cpu().detach().numpy()
        results = {
            'accuracy': accuracy_score(Y_test, Y_pred),
            'precision': precision_score(Y_test, Y_pred, average='weighted'),
            'recall': recall_score(Y_test, Y_pred, average='weighted'),
            'f1': f1_score(Y_test, Y_pred, average='weighted'),
            'confusion_matrix': confusion_matrix(Y_test, Y_pred),
            'classification_report': classification_report(Y_test, Y_pred)

```

```

    }
    self.last_results = results
    return results
def print_results(self, results=None):
    if results is None:
        try:
            results = self.last_results
        except:
            results = self.get_results()
    for key, value in results.items():
        if key in ['confusion_matrix', 'classification_report']:
            print(f'{key.capitalize()}: \n{value}')
        else:
            print(f'{key.capitalize()}: {value}')

class LogisticClassifier(Classifier):
    def __init__(self, input_dim=0, activation=nn.ReLU, hidden_layers = [64,
↪32, 16], pass_through=False):
        super().__init__()
        if pass_through:
            return
        self.stack_list = [nn.Linear(input_dim, hidden_layers[0]), activation()]
        for i in range(1, len(hidden_layers)):
            self.stack_list.extend([nn.Linear(hidden_layers[i-1],
↪hidden_layers[i]), activation()])

        self.stack_list.extend([nn.Linear(hidden_layers[-1], 1), nn.Sigmoid()])
        self.stack = nn.Sequential(*self.stack_list)

    def forward(self, x):
        return self.stack(x)

    def predict(self, x):
        with torch.no_grad():
            return self.forward(x).round()

    def train(self, epochs, X_train, X_test, Y_train, Y_test, alpha, loss_fn=nn.
↪BCELoss(), print_epoch=500):
        optimizer = torch.optim.SGD(self.parameters(), lr=alpha)

        for epoch in range(epochs):
            optimizer.zero_grad()
            Y_pred = self.forward(X_train)
            loss = loss_fn(Y_pred, Y_train)
            loss.backward()
            optimizer.step()

```

```

        if epoch % print_epoch == 0:
            test_loss = loss_fn(self.forward(X_test), Y_test)
            print(f'Epoch {epoch}: Training Loss: {loss.item()}, Test Loss: {
↪test_loss.item()}')
            Y_pred = self.predict(X_test)
            self.last_pred = Y_pred
            self.last_test = Y_test
            return [Y_test, Y_pred]

```

```

[69]: X_train_2 = torch.tensor(X_train_2).to(device).float()
X_test_2 = torch.tensor(X_test_2).to(device).float()
Y_train_2 = torch.tensor(Y_train_2).to(device).float().view(-1, 1)
Y_test_2 = torch.tensor(Y_test_2).to(device).float().view(-1, 1)

model_2a = LogisticClassifier(
    input_dim=X_train_2.shape[1],
    hidden_layers=[32],
    activation=nn.ReLU
).to(device)

model_2a.train(
    epochs=5000,
    X_train=X_train_2,
    X_test=X_test_2,
    Y_train=Y_train_2,
    Y_test=Y_test_2,
    alpha=1e-1,
    loss_fn=nn.BCELoss(),
    print_epoch=500
)

results_2a = model_2a.get_results()
model_2a.print_results(results_2a)

```

```

Epoch 0: Training Loss: 0.6775199174880981, Test Loss: 0.6508200168609619
Epoch 500: Training Loss: 0.05119786411523819, Test Loss: 0.08860453963279724
Epoch 1000: Training Loss: 0.03805161640048027, Test Loss: 0.09681770205497742
Epoch 1500: Training Loss: 0.029710467904806137, Test Loss: 0.10133590549230576
Epoch 2000: Training Loss: 0.023742537945508957, Test Loss: 0.09929131716489792
Epoch 2500: Training Loss: 0.018893906846642494, Test Loss: 0.09738191962242126
Epoch 3000: Training Loss: 0.015220258384943008, Test Loss: 0.09081994742155075
Epoch 3500: Training Loss: 0.012455631978809834, Test Loss: 0.08439692854881287
Epoch 4000: Training Loss: 0.010189911350607872, Test Loss: 0.07659681141376495
Epoch 4500: Training Loss: 0.008512535132467747, Test Loss: 0.06754373759031296
Accuracy: 0.9649122807017544
Precision: 0.9657164890247598
Recall: 0.9649122807017544

```


F1: 0.965011961722488

Confusion_matrix:

```
[[46  1]
 [ 3 64]]
```

Classification_report:

	precision	recall	f1-score	support
0.0	0.94	0.98	0.96	47
1.0	0.98	0.96	0.97	67
accuracy			0.96	114
macro avg	0.96	0.97	0.96	114
weighted avg	0.97	0.96	0.97	114

```
[70]: model_2b = LogisticClassifier(
        input_dim=X_train_2.shape[1],
        hidden_layers=[32,64,32],
        activation=nn.ReLU
    ).to(device)

    model_2b.train(
        epochs=5000,
        X_train=X_train_2,
        X_test=X_test_2,
        Y_train=Y_train_2,
        Y_test=Y_test_2,
        alpha=5e-2,
        loss_fn=nn.BCELoss(),
        print_epoch=500
    )
    print('\n')
    results_2b = model_2b.get_results()
    model_2b.print_results(results_2b)
```

```
Epoch 0: Training Loss: 0.7006031274795532, Test Loss: 0.6969746351242065
Epoch 500: Training Loss: 0.0444357804954052, Test Loss: 0.09696567803621292
Epoch 1000: Training Loss: 0.023338234052062035, Test Loss: 0.0965108647942543
Epoch 1500: Training Loss: 0.01155084278434515, Test Loss: 0.09876886010169983
Epoch 2000: Training Loss: 0.0062743364833295345, Test Loss: 0.1080765500664711
Epoch 2500: Training Loss: 0.0038169343024492264, Test Loss: 0.11713045835494995
Epoch 3000: Training Loss: 0.0025181507226079702, Test Loss: 0.12493595480918884
Epoch 3500: Training Loss: 0.0017743278294801712, Test Loss: 0.13145829737186432
Epoch 4000: Training Loss: 0.001317627727985382, Test Loss: 0.1369219720363617
Epoch 4500: Training Loss: 0.00102024688385427, Test Loss: 0.1416374146938324
```

Accuracy: 0.9473684210526315

Precision: 0.9502514456074828

Recall: 0.9473684210526315

F1: 0.9476328183095101

Confusion_matrix:

```
[[46  1]
 [ 5 62]]
```

Classification_report:

	precision	recall	f1-score	support
0.0	0.90	0.98	0.94	47
1.0	0.98	0.93	0.95	67
accuracy			0.95	114
macro avg	0.94	0.95	0.95	114
weighted avg	0.95	0.95	0.95	114

```
[71]: # Import cifar-10 dataset
from torchvision import datasets
import torchvision.transforms as transforms
from torch.utils.data.sampler import SubsetRandomSampler
from torch.utils.data import DataLoader

data_path = './data'
cifar10 = datasets.CIFAR10(data_path, train=True, download=True,
    ↪transform=transforms.ToTensor())

device = 'cuda' if torch.cuda.is_available() else 'cpu'
```

Files already downloaded and verified

```
[72]: train_imgs = torch.stack([img for img, _ in cifar10], dim=3)
view = train_imgs.view(3, -1)#.to(device=device)

mean = train_imgs.view(3, -1).mean(dim=1)
std = train_imgs.view(3, -1).std(dim=1)

transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(mean, std)
])

cifar10_train = datasets.CIFAR10(data_path, train=True, download=True,
    ↪transform=transform)
train_loader = DataLoader(cifar10_train, batch_size=64, shuffle=True)

cifar10_test = datasets.CIFAR10(data_path, train=False, download=True,
    ↪transform=transform)
```

```

test_loader = DataLoader(cifar10_test, batch_size=64, shuffle=False)
X_test_3 = torch.stack([img for img, _ in cifar10_test], dim=3).view(3, -1).
    ↳to(device=device)
Y_test_3 = torch.tensor([label for _, label in cifar10_test]).to(device=device)

```

Files already downloaded and verified

Files already downloaded and verified

```

[73]: import time

class ImageClassifier(Classifier):
    def __init__(self, input_dim=0, output_dim = 0, activation=nn.ReLU,
↳hidden_layers = [64, 32, 16], pass_through=False):
        super().__init__()
        self.stack_list = [nn.Flatten(), nn.Linear(input_dim,
↳hidden_layers[0]), activation()]
        for i in range(1, len(hidden_layers)):
            self.stack_list.extend([nn.Linear(hidden_layers[i-1],
↳hidden_layers[i]), activation()])

        self.stack_list.extend([nn.Linear(hidden_layers[-1], output_dim), nn.
↳Softmax(dim=1)])
        self.stack = nn.Sequential(*self.stack_list)
    def forward(self, x):
        return self.stack(x)
    def predict(self, x):
        with torch.no_grad():
            return self.forward(x).argmax(dim=1)
    def train_model(
        self,
        epochs,
        train_loader,
        test_loader,
        alpha,
        loss_fn=nn.CrossEntropyLoss(),
        optimizer=torch.optim.SGD,
        print_epoch=10,
    ):

        optimizer = optimizer(self.parameters(), lr=alpha)
        training_time = 0
        for epoch in range(epochs):
            self.train()

            start_time = time.time()
            train_loss = 0
            for X_batch, Y_batch in train_loader:

```

```

        X_batch, Y_batch = X_batch.to(device), Y_batch.to(device)
        optimizer.zero_grad()
        Y_pred = self.forward(X_batch)
        loss = loss_fn(Y_pred, Y_batch)
        loss.backward()
        optimizer.step()

        train_loss += loss.item()
        training_time += time.time() - start_time

    self.eval()
    with torch.no_grad():
        test_loss = 0
        Y_pred_eval = []
        Y_test = []
        for X_test_batch, Y_test_batch in test_loader:
            X_test_batch, Y_test_batch = X_test_batch.to(device),
↪Y_test_batch.to(device)

            out = self.forward(X_test_batch)
            test_loss += loss_fn(out, Y_test_batch).item()
            Y_test.extend(Y_test_batch.cpu().detach().numpy())
            Y_pred_eval.extend(out.argmax(dim=1).cpu().detach().numpy())

        accuracy = accuracy_score(Y_test, Y_pred_eval)
        if epoch % print_epoch == 0:
            print(f'Epoch {epoch}: Training Loss: {train_loss/
↪len(train_loader)}, Test Loss: {test_loss/len(test_loader)}, Accuracy:
↪{accuracy}')
            self.last_pred = torch.tensor(Y_pred_eval)
            self.last_test = torch.tensor(Y_test)
            print(f'\nTraining Time: {training_time} seconds\n')

```

```

[74]: device = 'cuda'
model_3a = ImageClassifier(
    input_dim=3*32*32,
    output_dim=10,
    hidden_layers=[256],
    activation=nn.Tanh
).to(device)

model_3a.train_model(
    epochs=100,

```

```

    train_loader=train_loader,
    test_loader=test_loader,
    alpha=1e-2,
    loss_fn=nn.CrossEntropyLoss(),
    print_epoch=1
)

model_3a.get_results()
model_3a.print_results()

```

```

Epoch 0: Training Loss: 2.201410284737492, Test Loss: 2.1488690786300952,
Accuracy: 0.3354
Epoch 1: Training Loss: 2.1313129562855986, Test Loss: 2.115134690217911,
Accuracy: 0.3624
Epoch 2: Training Loss: 2.1072356914315384, Test Loss: 2.097880730963057,
Accuracy: 0.3769
Epoch 3: Training Loss: 2.0923693166364488, Test Loss: 2.086663885481039,
Accuracy: 0.3865
Epoch 4: Training Loss: 2.08120296205706, Test Loss: 2.077612098614881,
Accuracy: 0.3979
Epoch 5: Training Loss: 2.0715074409609255, Test Loss: 2.070352053945991,
Accuracy: 0.4042
Epoch 6: Training Loss: 2.0632094161589736, Test Loss: 2.064615787973829,
Accuracy: 0.4078
Epoch 7: Training Loss: 2.0561082657340846, Test Loss: 2.0597093469777685,
Accuracy: 0.4105
Epoch 8: Training Loss: 2.0497846004298275, Test Loss: 2.054938104501955,
Accuracy: 0.4137
Epoch 9: Training Loss: 2.0441966700127057, Test Loss: 2.0527782402220804,
Accuracy: 0.415
Epoch 10: Training Loss: 2.0390897987748655, Test Loss: 2.048685533225916,
Accuracy: 0.4202
Epoch 11: Training Loss: 2.034578595319977, Test Loss: 2.0454648467385845,
Accuracy: 0.4231
Epoch 12: Training Loss: 2.0302584508190984, Test Loss: 2.0433450481694218,
Accuracy: 0.4246
Epoch 13: Training Loss: 2.02627748418647, Test Loss: 2.0410740231252777,
Accuracy: 0.4278
Epoch 14: Training Loss: 2.0223188958204616, Test Loss: 2.038747470090344,
Accuracy: 0.4294
Epoch 15: Training Loss: 2.018991274144643, Test Loss: 2.037671606252148,
Accuracy: 0.4305
Epoch 16: Training Loss: 2.0157555367635642, Test Loss: 2.035060773229903,
Accuracy: 0.4328
Epoch 17: Training Loss: 2.0125163214285964, Test Loss: 2.032249483333272,
Accuracy: 0.4377
Epoch 18: Training Loss: 2.0089872252300878, Test Loss: 2.0313054520613067,
Accuracy: 0.4373

```

Epoch 19: Training Loss: 2.006210855053514, Test Loss: 2.030170185550763,
Accuracy: 0.4373
Epoch 20: Training Loss: 2.0031930728031853, Test Loss: 2.0274508234801565,
Accuracy: 0.4399
Epoch 21: Training Loss: 2.000550609567891, Test Loss: 2.026248433787352,
Accuracy: 0.4396
Epoch 22: Training Loss: 1.9977973473956212, Test Loss: 2.02506713502726,
Accuracy: 0.441
Epoch 23: Training Loss: 1.9950800313973975, Test Loss: 2.0247168244829603,
Accuracy: 0.4403
Epoch 24: Training Loss: 1.9924540940453024, Test Loss: 2.0225466793509805,
Accuracy: 0.4421
Epoch 25: Training Loss: 1.990202148552136, Test Loss: 2.021040697006663,
Accuracy: 0.4442
Epoch 26: Training Loss: 1.9875995176832388, Test Loss: 2.02035410009372,
Accuracy: 0.4446
Epoch 27: Training Loss: 1.9850812461370093, Test Loss: 2.018848002336587,
Accuracy: 0.4461
Epoch 28: Training Loss: 1.982587840093676, Test Loss: 2.0185956916991312,
Accuracy: 0.4447
Epoch 29: Training Loss: 1.9802616491647023, Test Loss: 2.017457759304411,
Accuracy: 0.447
Epoch 30: Training Loss: 1.977693507585989, Test Loss: 2.0163778438689604,
Accuracy: 0.4484
Epoch 31: Training Loss: 1.975706407633584, Test Loss: 2.0154311193782055,
Accuracy: 0.4482
Epoch 32: Training Loss: 1.9732406907679174, Test Loss: 2.014418176025342,
Accuracy: 0.4508
Epoch 33: Training Loss: 1.9709903735029117, Test Loss: 2.0140721251250833,
Accuracy: 0.4491
Epoch 34: Training Loss: 1.968823590394481, Test Loss: 2.0130390232535684,
Accuracy: 0.4507
Epoch 35: Training Loss: 1.9664480809665397, Test Loss: 2.0122012309967334,
Accuracy: 0.4499
Epoch 36: Training Loss: 1.964585790853671, Test Loss: 2.0116470315653805,
Accuracy: 0.4507
Epoch 37: Training Loss: 1.9623820021024445, Test Loss: 2.0096353451917124,
Accuracy: 0.4548
Epoch 38: Training Loss: 1.9603788617931668, Test Loss: 2.0097902701918486,
Accuracy: 0.454
Epoch 39: Training Loss: 1.958145472704602, Test Loss: 2.0092169615873106,
Accuracy: 0.4545
Epoch 40: Training Loss: 1.9559663636300264, Test Loss: 2.0086839882431518,
Accuracy: 0.4548
Epoch 41: Training Loss: 1.9539883107785374, Test Loss: 2.0087792797453083,
Accuracy: 0.4545
Epoch 42: Training Loss: 1.9520726197820795, Test Loss: 2.007216927352225,
Accuracy: 0.4554

Epoch 43: Training Loss: 1.9498786005522588, Test Loss: 2.0065962007850597, Accuracy: 0.4567
Epoch 44: Training Loss: 1.9478924571705596, Test Loss: 2.0057929359423885, Accuracy: 0.4583
Epoch 45: Training Loss: 1.9459039570425478, Test Loss: 2.0053204077823907, Accuracy: 0.457
Epoch 46: Training Loss: 1.9441391784516746, Test Loss: 2.0043849876731823, Accuracy: 0.4595
Epoch 47: Training Loss: 1.9421423170572656, Test Loss: 2.0044837848396058, Accuracy: 0.4594
Epoch 48: Training Loss: 1.94016695022583, Test Loss: 2.002669891734032, Accuracy: 0.4619
Epoch 49: Training Loss: 1.9382048354429358, Test Loss: 2.002578695868231, Accuracy: 0.4605
Epoch 50: Training Loss: 1.9363141062924318, Test Loss: 2.0038897072433666, Accuracy: 0.4596
Epoch 51: Training Loss: 1.9344832413946575, Test Loss: 2.0018469095230103, Accuracy: 0.4612
Epoch 52: Training Loss: 1.9324918358832064, Test Loss: 2.0017270867232306, Accuracy: 0.4616
Epoch 53: Training Loss: 1.9307439822675017, Test Loss: 2.001353355729656, Accuracy: 0.4633
Epoch 54: Training Loss: 1.928962046685426, Test Loss: 2.0016647691179994, Accuracy: 0.4608
Epoch 55: Training Loss: 1.927043501068564, Test Loss: 1.9999224835899985, Accuracy: 0.4651
Epoch 56: Training Loss: 1.9250040407985678, Test Loss: 1.9994656105709683, Accuracy: 0.4641
Epoch 57: Training Loss: 1.9234034970898153, Test Loss: 1.9991518092003597, Accuracy: 0.4642
Epoch 58: Training Loss: 1.9215860725058924, Test Loss: 1.999201297000715, Accuracy: 0.4643
Epoch 59: Training Loss: 1.9198868384446635, Test Loss: 1.9987571254657333, Accuracy: 0.4647
Epoch 60: Training Loss: 1.9179791044396208, Test Loss: 1.9974228773906733, Accuracy: 0.4659
Epoch 61: Training Loss: 1.916289828469991, Test Loss: 1.9975098523364705, Accuracy: 0.4673
Epoch 62: Training Loss: 1.9146924146910762, Test Loss: 1.9978118755255536, Accuracy: 0.4662
Epoch 63: Training Loss: 1.912617788137987, Test Loss: 1.9964036083525154, Accuracy: 0.4676
Epoch 64: Training Loss: 1.9110569766415355, Test Loss: 1.995940009499811, Accuracy: 0.468
Epoch 65: Training Loss: 1.9095145576750225, Test Loss: 1.9957442184922043, Accuracy: 0.4668
Epoch 66: Training Loss: 1.9077311972218096, Test Loss: 1.995925371813926, Accuracy: 0.469

Epoch 67: Training Loss: 1.906208941698684, Test Loss: 1.995034005231918,
Accuracy: 0.4691
Epoch 68: Training Loss: 1.9042365439712543, Test Loss: 1.9958832036158083,
Accuracy: 0.4686
Epoch 69: Training Loss: 1.9026686892180187, Test Loss: 1.994063199705379,
Accuracy: 0.4693
Epoch 70: Training Loss: 1.9011069923410635, Test Loss: 1.9941022388494698,
Accuracy: 0.4689
Epoch 71: Training Loss: 1.89944163963313, Test Loss: 1.9938461734990405,
Accuracy: 0.4708
Epoch 72: Training Loss: 1.8979422139084858, Test Loss: 1.9931916909612668,
Accuracy: 0.4695
Epoch 73: Training Loss: 1.8962264480188376, Test Loss: 1.9930588234761717,
Accuracy: 0.4729
Epoch 74: Training Loss: 1.894419810503645, Test Loss: 1.9929505951085669,
Accuracy: 0.4722
Epoch 75: Training Loss: 1.8932654946051595, Test Loss: 1.9926604366606209,
Accuracy: 0.469
Epoch 76: Training Loss: 1.8913014256740774, Test Loss: 1.9919216412647514,
Accuracy: 0.4722
Epoch 77: Training Loss: 1.8898866644600774, Test Loss: 1.9922871020189516,
Accuracy: 0.4712
Epoch 78: Training Loss: 1.8885408505759276, Test Loss: 1.9918207097205387,
Accuracy: 0.4715
Epoch 79: Training Loss: 1.8867450173553604, Test Loss: 1.9917527133492148,
Accuracy: 0.4733
Epoch 80: Training Loss: 1.885162004424483, Test Loss: 1.9901161619052765,
Accuracy: 0.4739
Epoch 81: Training Loss: 1.8836467884995443, Test Loss: 1.990831255153486,
Accuracy: 0.4749
Epoch 82: Training Loss: 1.8821341834409768, Test Loss: 1.9903806865594948,
Accuracy: 0.4735
Epoch 83: Training Loss: 1.880492230053143, Test Loss: 1.9894001423173648,
Accuracy: 0.4761
Epoch 84: Training Loss: 1.8792056093740341, Test Loss: 1.9900260746099387,
Accuracy: 0.474
Epoch 85: Training Loss: 1.877844167030071, Test Loss: 1.9894564698456199,
Accuracy: 0.4731
Epoch 86: Training Loss: 1.876429661765428, Test Loss: 1.9892730614182297,
Accuracy: 0.4771
Epoch 87: Training Loss: 1.8750434484323273, Test Loss: 1.9887491677217424,
Accuracy: 0.476
Epoch 88: Training Loss: 1.8735629942106165, Test Loss: 1.9892685428546493,
Accuracy: 0.4746
Epoch 89: Training Loss: 1.8719964249969443, Test Loss: 1.9882509078189825,
Accuracy: 0.4746
Epoch 90: Training Loss: 1.8706222033256765, Test Loss: 1.987908758175601,
Accuracy: 0.4747

Epoch 91: Training Loss: 1.869113135063435, Test Loss: 1.9885168736148033,
 Accuracy: 0.476
 Epoch 92: Training Loss: 1.8677857948081267, Test Loss: 1.988387291598472,
 Accuracy: 0.4748
 Epoch 93: Training Loss: 1.8664771082151272, Test Loss: 1.9878002739256355,
 Accuracy: 0.475
 Epoch 94: Training Loss: 1.8650445026509903, Test Loss: 1.9875242808821854,
 Accuracy: 0.4747
 Epoch 95: Training Loss: 1.8636897526433707, Test Loss: 1.9874323697606469,
 Accuracy: 0.475
 Epoch 96: Training Loss: 1.8621286147695673, Test Loss: 1.9872268529454613,
 Accuracy: 0.4757
 Epoch 97: Training Loss: 1.8607349558864408, Test Loss: 1.9873499710848377,
 Accuracy: 0.4763
 Epoch 98: Training Loss: 1.859359428248442, Test Loss: 1.986620311524458,
 Accuracy: 0.4745
 Epoch 99: Training Loss: 1.8583002029477482, Test Loss: 1.986218580015146,
 Accuracy: 0.4742

Training Time: 1157.4642927646637 seconds

Accuracy: 0.4742

Precision: 0.46611456870171314

Recall: 0.4742

F1: 0.4682781879764696

Confusion_matrix:

```
[[587  37  61  25  20  17  24  36 137  56]
 [ 34 561  26  39  21  23  41  38  77 140]
 [105  37 289  78 141  69 143  69  42  27]
 [ 37  37  93 271  65 163 149  61  42  82]
 [ 50  25 128  58 394  57 116 111  31  30]
 [ 29  28 100 156  80 348  92  84  50  33]
 [ 11  24  58  70 127  63 563  29  25  30]
 [ 45  37  54  50  85  76  42 518  25  68]
 [ 89  58  16  24  13  28  14  20 681  57]
 [ 64 158  15  28  16  26  49  42  72 530]]
```

Classification_report:

	precision	recall	f1-score	support
0	0.56	0.59	0.57	1000
1	0.56	0.56	0.56	1000
2	0.34	0.29	0.31	1000
3	0.34	0.27	0.30	1000
4	0.41	0.39	0.40	1000
5	0.40	0.35	0.37	1000
6	0.46	0.56	0.50	1000
7	0.51	0.52	0.52	1000
8	0.58	0.68	0.62	1000

	9	0.50	0.53	0.52	1000
accuracy				0.47	10000
macro avg		0.47	0.47	0.47	10000
weighted avg		0.47	0.47	0.47	10000

```
[75]: model_3a_results = model_3a.get_results()
      model_3a.print_results(model_3a_results)
```

```
Accuracy: 0.4742
Precision: 0.46611456870171314
Recall: 0.4742
F1: 0.4682781879764696
Confusion_matrix:
[[587  37  61  25  20  17  24  36 137  56]
 [ 34 561  26  39  21  23  41  38  77 140]
 [105  37 289  78 141  69 143  69  42  27]
 [ 37  37  93 271  65 163 149  61  42  82]
 [ 50  25 128  58 394  57 116 111  31  30]
 [ 29  28 100 156  80 348  92  84  50  33]
 [ 11  24  58  70 127  63 563  29  25  30]
 [ 45  37  54  50  85  76  42 518  25  68]
 [ 89  58  16  24  13  28  14  20 681  57]
 [ 64 158  15  28  16  26  49  42  72 530]]
Classification_report:
      precision    recall  f1-score   support

0         0.56      0.59      0.57       1000
1         0.56      0.56      0.56       1000
2         0.34      0.29      0.31       1000
3         0.34      0.27      0.30       1000
4         0.41      0.39      0.40       1000
5         0.40      0.35      0.37       1000
6         0.46      0.56      0.50       1000
7         0.51      0.52      0.52       1000
8         0.58      0.68      0.62       1000
9         0.50      0.53      0.52       1000

accuracy
macro avg
weighted avg
```

				0.47	10000
		0.47	0.47	0.47	10000
		0.47	0.47	0.47	10000

Noticed the SGD optimizer is much slower than Adam, but yields more consistent results. To improve the model performance, the alpha was increased to 1e-2 to deal with the lower learning rate of SGD

```
[76]: device = 'cuda'
model_3b = ImageClassifier(
    input_dim=3*32*32,
    output_dim=10,
    hidden_layers=[256,512,128],#[256,384,256],
    activation=nn.Tanh
).to(device)

model_3b.train_model(
    epochs=100,
    train_loader=train_loader,
    test_loader=test_loader,
    alpha=1e-2,
    loss_fn=nn.CrossEntropyLoss(),
    optimizer = torch.optim.SGD,
    print_epoch=1
)

model_3b_results = model_3a.get_results()
model_3b.print_results(model_3a_results)
```

```
Epoch 0: Training Loss: 2.280403497273965, Test Loss: 2.252564680804113,
Accuracy: 0.2359
Epoch 1: Training Loss: 2.228228701045141, Test Loss: 2.202385954036834,
Accuracy: 0.2707
Epoch 2: Training Loss: 2.1879612088508313, Test Loss: 2.172818032039958,
Accuracy: 0.2915
Epoch 3: Training Loss: 2.164266648804745, Test Loss: 2.152656055559778,
Accuracy: 0.316
Epoch 4: Training Loss: 2.145001944983402, Test Loss: 2.1356226805668728,
Accuracy: 0.3354
Epoch 5: Training Loss: 2.128756971615355, Test Loss: 2.1195827631434057,
Accuracy: 0.3519
Epoch 6: Training Loss: 2.113841903484081, Test Loss: 2.107796171668229,
Accuracy: 0.3608
Epoch 7: Training Loss: 2.1028579657949753, Test Loss: 2.098387276291088,
Accuracy: 0.371
Epoch 8: Training Loss: 2.0934950303848443, Test Loss: 2.090748143803542,
Accuracy: 0.3772
Epoch 9: Training Loss: 2.0848870728631765, Test Loss: 2.084034829382684,
Accuracy: 0.3808
Epoch 10: Training Loss: 2.0770044459406374, Test Loss: 2.078609412642801,
Accuracy: 0.3871
Epoch 11: Training Loss: 2.070014102532126, Test Loss: 2.0740907898374425,
Accuracy: 0.391
Epoch 12: Training Loss: 2.0636268466939707, Test Loss: 2.0721085997903423,
Accuracy: 0.3919
Epoch 13: Training Loss: 2.0574979936070457, Test Loss: 2.0637668682511445,
```

Accuracy: 0.3985
Epoch 14: Training Loss: 2.051519621058803, Test Loss: 2.0614985918543143, Accuracy: 0.4013
Epoch 15: Training Loss: 2.046250664821976, Test Loss: 2.056406200311746, Accuracy: 0.4054
Epoch 16: Training Loss: 2.0412039791836456, Test Loss: 2.053648417163047, Accuracy: 0.4074
Epoch 17: Training Loss: 2.036012465539186, Test Loss: 2.049658063111032, Accuracy: 0.4128
Epoch 18: Training Loss: 2.031437578561056, Test Loss: 2.0488287567333052, Accuracy: 0.4131
Epoch 19: Training Loss: 2.0267347254411643, Test Loss: 2.0439856819286466, Accuracy: 0.4178
Epoch 20: Training Loss: 2.022114611952506, Test Loss: 2.043041490445471, Accuracy: 0.4191
Epoch 21: Training Loss: 2.017402906704437, Test Loss: 2.038575057011501, Accuracy: 0.4234
Epoch 22: Training Loss: 2.0128561633322244, Test Loss: 2.0373852860396076, Accuracy: 0.4235
Epoch 23: Training Loss: 2.0082216503675028, Test Loss: 2.0335943714068954, Accuracy: 0.4275
Epoch 24: Training Loss: 2.0037876928553864, Test Loss: 2.030826511656403, Accuracy: 0.4319
Epoch 25: Training Loss: 1.9994239329986865, Test Loss: 2.0289009050199183, Accuracy: 0.4343
Epoch 26: Training Loss: 1.9945184483247644, Test Loss: 2.0266629237278253, Accuracy: 0.4382
Epoch 27: Training Loss: 1.9897202751825533, Test Loss: 2.0271820012171555, Accuracy: 0.437
Epoch 28: Training Loss: 1.9846469432191776, Test Loss: 2.0247295221705346, Accuracy: 0.4392
Epoch 29: Training Loss: 1.9798921794842577, Test Loss: 2.0211990319999162, Accuracy: 0.4406
Epoch 30: Training Loss: 1.975139911522341, Test Loss: 2.019865758859428, Accuracy: 0.4439
Epoch 31: Training Loss: 1.970379686569009, Test Loss: 2.0177729889086096, Accuracy: 0.4455
Epoch 32: Training Loss: 1.9653151279215313, Test Loss: 2.015972366758213, Accuracy: 0.446
Epoch 33: Training Loss: 1.9602141008352685, Test Loss: 2.0135348678394487, Accuracy: 0.4484
Epoch 34: Training Loss: 1.9551919712434949, Test Loss: 2.011253961332285, Accuracy: 0.4515
Epoch 35: Training Loss: 1.950315807481556, Test Loss: 2.0147219919095374, Accuracy: 0.4458
Epoch 36: Training Loss: 1.9453664797041423, Test Loss: 2.0082669774438164, Accuracy: 0.4536
Epoch 37: Training Loss: 1.940512748782897, Test Loss: 2.009062681987787,

Accuracy: 0.4502
Epoch 38: Training Loss: 1.9353905344558189, Test Loss: 2.009172895152098, Accuracy: 0.4531
Epoch 39: Training Loss: 1.9309233869128215, Test Loss: 2.0051034256151525, Accuracy: 0.457
Epoch 40: Training Loss: 1.9258459394850085, Test Loss: 2.0054075057339515, Accuracy: 0.4542
Epoch 41: Training Loss: 1.920948523389714, Test Loss: 2.0034787070219684, Accuracy: 0.4572
Epoch 42: Training Loss: 1.9158197467589317, Test Loss: 2.003737722232843, Accuracy: 0.4566
Epoch 43: Training Loss: 1.9112602746700083, Test Loss: 2.007011801573881, Accuracy: 0.45
Epoch 44: Training Loss: 1.9068707323745084, Test Loss: 2.003478769284145, Accuracy: 0.4543
Epoch 45: Training Loss: 1.90230433227461, Test Loss: 2.00412516609119, Accuracy: 0.4558
Epoch 46: Training Loss: 1.8979026368816796, Test Loss: 2.0014274537942973, Accuracy: 0.4565
Epoch 47: Training Loss: 1.8929194580868383, Test Loss: 2.005838249139725, Accuracy: 0.4531
Epoch 48: Training Loss: 1.8888147938281983, Test Loss: 2.0037533644657985, Accuracy: 0.4546
Epoch 49: Training Loss: 1.8840698384872787, Test Loss: 2.0042982344414777, Accuracy: 0.4539
Epoch 50: Training Loss: 1.8803379241462863, Test Loss: 2.0060581659815115, Accuracy: 0.4536
Epoch 51: Training Loss: 1.8755995143405007, Test Loss: 2.0026044678536192, Accuracy: 0.4558
Epoch 52: Training Loss: 1.8712224228607723, Test Loss: 1.9998929014631137, Accuracy: 0.4587
Epoch 53: Training Loss: 1.8668004260648547, Test Loss: 2.0050053262406853, Accuracy: 0.4537
Epoch 54: Training Loss: 1.862304040080751, Test Loss: 1.9967357801024321, Accuracy: 0.4616
Epoch 55: Training Loss: 1.8581476738995604, Test Loss: 2.011644670158435, Accuracy: 0.4476
Epoch 56: Training Loss: 1.85449959600673, Test Loss: 2.0013837252452875, Accuracy: 0.4555
Epoch 57: Training Loss: 1.8503952450154688, Test Loss: 2.005105645793259, Accuracy: 0.4534
Epoch 58: Training Loss: 1.8461702011735237, Test Loss: 2.0121096296674885, Accuracy: 0.4455
Epoch 59: Training Loss: 1.8425086915035687, Test Loss: 2.00076228418168, Accuracy: 0.4561
Epoch 60: Training Loss: 1.8384573668470163, Test Loss: 1.9964367902962266, Accuracy: 0.4612
Epoch 61: Training Loss: 1.834525834568931, Test Loss: 1.9943766305401067,

Accuracy: 0.4647
Epoch 62: Training Loss: 1.8304467589958855, Test Loss: 1.9952366450789627, Accuracy: 0.4618
Epoch 63: Training Loss: 1.8265629633308371, Test Loss: 2.0005740686586706, Accuracy: 0.456
Epoch 64: Training Loss: 1.822805018223765, Test Loss: 2.005883834923908, Accuracy: 0.4491
Epoch 65: Training Loss: 1.81961776914499, Test Loss: 1.99604929632442, Accuracy: 0.4605
Epoch 66: Training Loss: 1.8148655074331768, Test Loss: 1.9966157560895204, Accuracy: 0.4607
Epoch 67: Training Loss: 1.8123863682417614, Test Loss: 2.001093756621051, Accuracy: 0.4567
Epoch 68: Training Loss: 1.8089882707047036, Test Loss: 1.9976130207632756, Accuracy: 0.4591
Epoch 69: Training Loss: 1.8050172597246097, Test Loss: 2.000473508409634, Accuracy: 0.4585
Epoch 70: Training Loss: 1.8015150918680078, Test Loss: 1.9935544228098194, Accuracy: 0.4642
Epoch 71: Training Loss: 1.798106653458627, Test Loss: 1.9967582954722605, Accuracy: 0.459
Epoch 72: Training Loss: 1.7946201843373917, Test Loss: 1.9966439912273626, Accuracy: 0.4595
Epoch 73: Training Loss: 1.7911296585941558, Test Loss: 1.9975928455401377, Accuracy: 0.4605
Epoch 74: Training Loss: 1.7893220724352181, Test Loss: 1.9976800953506664, Accuracy: 0.4607
Epoch 75: Training Loss: 1.7857283984913546, Test Loss: 1.997754140264669, Accuracy: 0.461
Epoch 76: Training Loss: 1.7820604012140533, Test Loss: 2.0081516086675557, Accuracy: 0.4476
Epoch 77: Training Loss: 1.7795437854879044, Test Loss: 2.014743470841912, Accuracy: 0.442
Epoch 78: Training Loss: 1.7772340364468373, Test Loss: 1.9949941035288914, Accuracy: 0.4622
Epoch 79: Training Loss: 1.774098057423711, Test Loss: 2.003993206722721, Accuracy: 0.4537
Epoch 80: Training Loss: 1.771172871979911, Test Loss: 1.996228008513238, Accuracy: 0.4629
Epoch 81: Training Loss: 1.7697991905614847, Test Loss: 1.998249265039043, Accuracy: 0.4581
Epoch 82: Training Loss: 1.7668739221894834, Test Loss: 2.0033721316392255, Accuracy: 0.4551
Epoch 83: Training Loss: 1.7646816086281292, Test Loss: 1.9957349505394129, Accuracy: 0.4613
Epoch 84: Training Loss: 1.7610802517827515, Test Loss: 1.9970971710363012, Accuracy: 0.4589
Epoch 85: Training Loss: 1.7594040022481738, Test Loss: 1.995714901359218,

Accuracy: 0.4591
 Epoch 86: Training Loss: 1.7561884853236205, Test Loss: 2.000106747742671,
 Accuracy: 0.4548
 Epoch 87: Training Loss: 1.7541905739118375, Test Loss: 1.9980947082969034,
 Accuracy: 0.4594
 Epoch 88: Training Loss: 1.752043614302145, Test Loss: 1.9978720876061993,
 Accuracy: 0.4585
 Epoch 89: Training Loss: 1.7497063177016081, Test Loss: 2.04112223758819,
 Accuracy: 0.4134
 Epoch 90: Training Loss: 1.7477915985200105, Test Loss: 1.9977301518628552,
 Accuracy: 0.4571
 Epoch 91: Training Loss: 1.745197372515793, Test Loss: 2.0041657192691877,
 Accuracy: 0.4528
 Epoch 92: Training Loss: 1.7442166652825788, Test Loss: 1.998430759284147,
 Accuracy: 0.4579
 Epoch 93: Training Loss: 1.7417706339560506, Test Loss: 2.000380478087504,
 Accuracy: 0.4567
 Epoch 94: Training Loss: 1.7388503150561887, Test Loss: 2.007878496388721,
 Accuracy: 0.4491
 Epoch 95: Training Loss: 1.7377261348697535, Test Loss: 2.003134667493735,
 Accuracy: 0.4536
 Epoch 96: Training Loss: 1.7351144726014198, Test Loss: 2.000657352672261,
 Accuracy: 0.4562
 Epoch 97: Training Loss: 1.7338576975380977, Test Loss: 2.0077067985656156,
 Accuracy: 0.4488
 Epoch 98: Training Loss: 1.7315342722036648, Test Loss: 2.003403690969868,
 Accuracy: 0.4531
 Epoch 99: Training Loss: 1.730216651621377, Test Loss: 2.002178602917179,
 Accuracy: 0.4532

Training Time: 1186.6000225543976 seconds

Accuracy: 0.4742
 Precision: 0.46611456870171314
 Recall: 0.4742
 F1: 0.4682781879764696

Confusion_matrix:

```

[[587  37  61  25  20  17  24  36 137  56]
 [ 34 561  26  39  21  23  41  38  77 140]
 [105  37 289  78 141  69 143  69  42  27]
 [ 37  37  93 271  65 163 149  61  42  82]
 [ 50  25 128  58 394  57 116 111  31  30]
 [ 29  28 100 156  80 348  92  84  50  33]
 [ 11  24  58  70 127  63 563  29  25  30]
 [ 45  37  54  50  85  76  42 518  25  68]
 [ 89  58  16  24  13  28  14  20 681  57]
 [ 64 158  15  28  16  26  49  42  72 530]]
  
```

Classification_report:

	precision	recall	f1-score	support
0	0.56	0.59	0.57	1000
1	0.56	0.56	0.56	1000
2	0.34	0.29	0.31	1000
3	0.34	0.27	0.30	1000
4	0.41	0.39	0.40	1000
5	0.40	0.35	0.37	1000
6	0.46	0.56	0.50	1000
7	0.51	0.52	0.52	1000
8	0.58	0.68	0.62	1000
9	0.50	0.53	0.52	1000
accuracy			0.47	10000
macro avg	0.47	0.47	0.47	10000
weighted avg	0.47	0.47	0.47	10000