

Kabir__Notebook__7

December 9, 2024

[]:

```
[1]: import torch  
from torch import nn  
import torch.nn.functional as F  
from collections import OrderedDict  
from sklearn.metrics import accuracy_score, precision_score, recall_score,  
    f1_score, confusion_matrix, classification_report  
import matplotlib.pyplot as plt  
from torchvision import datasets  
import torchvision.transforms as transforms  
from torch.utils.data import DataLoader  
import time  
from IPython.core.magic import register_cell_magic  
import gc  
from torch.autograd import autograd, GradScaler  
from torch.utils.data import CudaDataPrefetcher  
  
@register_cell_magic  
def skip(line, cell):  
    return  
  
device = 'cuda'
```

```
/home/super/.local/lib/python3.11/site-packages/tqdm/auto.py:21: TqdmWarning:
IPProgress not found. Please update jupyter and ipywidgets. See
https://ipywidgets.readthedocs.io/en/stable/user_install.html
  from .autonotebook import tqdm as notebook tqdm
```

```
[2]: dl = False
data_path = './data'

# Load CIFAR-10 dataset with the simple transform
cifar10_train = datasets.CIFAR10(data_path, train=True, download=dl,
    ↪ transform=transforms.ToTensor())
try:
```

```

mean = torch.load('data/mean.pt')
std = torch.load('data/std.pt')
except FileNotFoundError:
    print("Computing Mean and Std")
    train_imgs = torch.stack([img for img, _ in cifar10_train], dim=3)#.
    ↪to(device=device)
    view = train_imgs.view(3, -1)#.to(device=device)

    mean = train_imgs.view(3, -1).mean(dim=1)
    std = train_imgs.view(3, -1).std(dim=1)

    transform = transforms.Compose([
        transforms.ToTensor(),
        transforms.Normalize(mean, std)
    ])

    torch.save(mean, 'data/mean.pt')
    torch.save(std, 'data/std.pt')

# Define the transform with normalization
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(mean, std)
])
print("Mean: ", mean)
print("Std: ", std)
cifar10_train = datasets.CIFAR10(data_path, train=True, download=dl,
    ↪transform=transform)
cifar10_test = datasets.CIFAR10(data_path, train=False, download=dl,
    ↪transform=transform)

```

/tmp/ipykernel_1721992/2313373535.py:7: FutureWarning: You are using `torch.load` with `weights_only=False` (the current default value), which uses the default pickle module implicitly. It is possible to construct malicious pickle data which will execute arbitrary code during unpickling (See <https://github.com/pytorch/pytorch/blob/main/SECURITY.md#untrusted-models> for more details). In a future release, the default value for `weights_only` will be flipped to `True`. This limits the functions that could be executed during unpickling. Arbitrary objects will no longer be allowed to be loaded via this mode unless they are explicitly allowlisted by the user via `torch.serialization.add_safe_globals`. We recommend you start setting `weights_only=True` for any use case where you don't have full control of the loaded file. Please open an issue on GitHub for any issues related to this experimental feature.

```

mean = torch.load('data/mean.pt')
/tmp/ipykernel_1721992/2313373535.py:8: FutureWarning: You are using
`torch.load` with `weights_only=False` (the current default value), which uses

```

the default pickle module implicitly. It is possible to construct malicious pickle data which will execute arbitrary code during unpickling (See <https://github.com/pytorch/pytorch/blob/main/SECURITY.md#untrusted-models> for more details). In a future release, the default value for `weights_only` will be flipped to `True`. This limits the functions that could be executed during unpickling. Arbitrary objects will no longer be allowed to be loaded via this mode unless they are explicitly allowlisted by the user via `torch.serialization.add_safe_globals`. We recommend you start setting `weights_only=True` for any use case where you don't have full control of the loaded file. Please open an issue on GitHub for any issues related to this experimental feature.

```
std = torch.load('data/std.pt')
```

```
Mean: tensor([0.4914, 0.4822, 0.4465])
```

```
Std: tensor([0.2470, 0.2435, 0.2616])
```

```
[3]: class Classifier(nn.Module):
    @classmethod
    def compare_results(cls, results1, results2):
        print('Comparing results:')
        comparisons = {
            'accuracy': 100*(results1['accuracy'] - results2['accuracy'])/
↪results1['accuracy'],
            'precision': 100*(results1['precision'] - results2['precision'])/
↪results1['precision'],
            'recall': 100*(results1['recall'] - results2['recall'])/
↪results1['recall'],
            'f1': 100*(results1['f1'] - results2['f1'])/results1['f1']
        }
        for key, value in comparisons.items():
            print(f'{key}: {value} %')

    def __init__(self):
        super().__init__()

    def forward(self, x):
        return self.stack(x)

    def predict(self, x):
        with torch.no_grad():
            self.eval()
            return self.forward(x).argmax(dim=1)

    def train_model(
        self,
        epochs,
        train_loader,
        test_loader,
        train_len,
```

```

test_len,
test_size,
loss_fn=nn.CrossEntropyLoss(),
optimizer=torch.optim.SGD,
optimizer_args = [],
optimizer_kwargs = {},
print_epoch=10,
header_epoch = 15,
sched_factor = 0.1,
sched_patience = 5
):

    scaler = GradScaler("cuda")
    optimizer = optimizer(self.parameters(), *optimizer_args,
↪**optimizer_kwargs)
    scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer,
↪'max', patience=sched_patience, factor=sched_factor)
    training_time = 0
    train_hist = torch.zeros(epochs, device=device)
    test_hist = torch.zeros(epochs, device=device)
    accuracy_hist = torch.zeros(epochs, device=device)

    cell_width = 20
    header_form_spec = f'^{cell_width}'

    epoch_inspection = {
        "Epoch": 0,
        "Epoch Time (s)": 0,
        "Training Loss": 0,
        "Test Loss ": 0,
        "Overfit (%)": 0,
        "Accuracy (%)": 0,
        "Δ Accuracy (%)": 0,
        "GPU Memory (GiB)": 0
    }

    header_string = "|"
    for key in epoch_inspection.keys():
        header_string += (f"{key:{header_form_spec}}|")

    divider_string = '-'*len(header_string)
    if print_epoch:
        print(f'Training {self.__class__.__name__}\n')
        print(divider_string)
    max_accuracy = torch.zeros(1, device=device)
    for epoch in range(epochs):
        begin_epoch = time.time()

```

```

self.train()

start_time = time.time()
train_loss = 0
for X_batch, Y_batch in train_loader:
    #X_batch, Y_batch = X_batch.to(device, non_blocking=True),
    ↪Y_batch.to(device, non_blocking=True)
    optimizer.zero_grad(set_to_none=True)
    with autocast("cuda"):
        Y_pred = self.forward(X_batch)
        loss = loss_fn(Y_pred, Y_batch)
        scaler.scale(loss).backward()
        scaler.step(optimizer)
        scaler.update()

    train_loss += loss
training_time += time.time() - start_time

train_loss = train_loss/train_len
train_hist[epoch] = train_loss

self.eval()
with torch.no_grad():
    test_loss = torch.zeros(1, device=device)
    correct = torch.zeros(1, device=device)

    for X_test_batch, Y_test_batch in test_loader:
        #X_test_batch, Y_test_batch = X_test_batch.to(device,
        ↪non_blocking=True), Y_test_batch.to(device, non_blocking=True)

        out = self.forward(X_test_batch)
        test_loss += loss_fn(out, Y_test_batch)
        correct += (out.argmax(dim=1) == Y_test_batch).sum()

test_loss = test_loss/test_len
test_hist[epoch] = test_loss
accuracy = correct/test_size
accuracy_hist[epoch] = accuracy

scheduler.step(accuracy)

end_epoch = time.time()
if print_epoch and (epoch % print_epoch == 0 or epoch == epochs -
↪1) :
    mem = (torch.cuda.memory_allocated() + torch.cuda.
    ↪memory_reserved())/1024**3

```

```

        if header_epoch and epoch % header_epoch == 0:
            print(header_string)
            print(divider_string)
            epoch_duration = end_epoch - begin_epoch
            overfit = 100 * (test_loss - train_loss) / train_loss
            d_accuracy = torch.zeros(1) if max_accuracy == 0 else 100 * ↵
↵(accuracy - max_accuracy) / max_accuracy
            if accuracy > max_accuracy:
                max_accuracy = accuracy

            epoch_inspection['Epoch'] = f'{epoch}'
            epoch_inspection['Epoch Time (s)'] = f'{epoch_duration:4f}'
            epoch_inspection['Training Loss'] = f'{train_loss.item():8f}'
            epoch_inspection['Test Loss '] = f'{test_loss.item():8f}'
            epoch_inspection['Overfit (%)'] = f'{overfit.item():4f}'
            epoch_inspection['Accuracy (%)'] = f'{accuracy.item()*100:4f}'
            epoch_inspection['Δ Accuracy (%)'] = f'{d_accuracy.item():4f}'
            epoch_inspection["GPU Memory (GiB)"] = f'{mem:2f}'
            for value in epoch_inspection.values():
                print(f"|{value:~{cell_width}}", end='')
            print('|')
            print(divider_string)

    print(f'\nTraining Time: {training_time} seconds\n')

    self.train_hist = train_hist
    self.test_hist = test_hist
    self.accuracy_hist = accuracy_hist

    def plot_training(self, title='Training Results'):
        plt.plot(self.train_hist.detach().cpu(), label='Training Loss')
        plt.plot(self.test_hist.detach().cpu(), label='Test Loss')
        plt.plot(self.accuracy_hist.detach().cpu(), label='Accuracy')
        plt.title(title)
        plt.xlabel('Epoch')
        plt.ylabel('Loss')
        plt.legend()
        plt.show()

    def get_results(self, Y_test=None, Y_pred=None):
        if Y_test is None:
            Y_test = self.last_test
        if Y_pred is None:
            Y_pred = self.last_pred

        if isinstance(Y_test, torch.Tensor):

```

```

        Y_test = Y_test.cpu().detach().numpy()
    if isinstance(Y_pred, torch.Tensor):
        Y_pred = Y_pred.cpu().detach().numpy()
    results = {
        'accuracy': accuracy_score(Y_test, Y_pred),
        'precision': precision_score(Y_test, Y_pred, average='weighted'),
        'recall': recall_score(Y_test, Y_pred, average='weighted'),
        'f1': f1_score(Y_test, Y_pred, average='weighted'),
        'confusion_matrix': confusion_matrix(Y_test, Y_pred),
        'classification_report': classification_report(Y_test, Y_pred)
    }
    self.last_results = results
    return results
def print_results(self, results=None):
    if results is None:
        try:
            results = self.last_results
        except:
            results = self.get_results()
    for key, value in results.items():
        if key in ['confusion_matrix', 'classification_report']:
            print(f'{key.capitalize()}: \n{value}')
        else:
            print(f'{key.capitalize()}: {value}')

```

```

[4]: class ConvImageClassifier(Classifier):
    def __init__(self, input_dim, conv_layers, fc_layers, activation=nn.ReLU):
        super().__init__()

        self.stack = nn.Sequential(OrderedDict(
            [
                ('conv0', nn.Conv2d(in_channels=3, out_channels=conv_layers[0],
↪kernel_size=3, padding=1)),
                ('activation0', activation()),
                ('maxpool0', nn.MaxPool2d(2)),
            ]
        ))

        for i in range(1, len(conv_layers)):
            self.stack.add_module(f'conv{i}', nn.
↪Conv2d(in_channels=conv_layers[i-1], out_channels=conv_layers[i],
↪kernel_size=3, padding=1))
            self.stack.add_module(f'activation{i}', activation())
            self.stack.add_module(f'maxpool{i}', nn.MaxPool2d(2))

        conv_out = input_dim // (2**len(conv_layers))
        self.stack.add_module('flatten', nn.Flatten())

```

```

        self.stack.add_module(f'fc0', nn.Linear(conv_out**2*conv_layers[-1],
↪fc_layers[0]))

        for i in range(1, len(fc_layers)):
            self.stack.add_module(f'activation_fc{i}', nn.Tanh())
            self.stack.add_module(f'fc{i}', nn.Linear(fc_layers[i-1],
↪fc_layers[i]))

```

```

[5]: try:
    del train_loader
    del test_loader
    del model_1a
    del model_1b
    del resnet
    del train_loader_cuda
    del test_loader_cuda
except:
    pass

# Reset CUDA context
start = time.time()
torch.cuda.empty_cache()
torch.cuda.reset_peak_memory_stats()

gc.collect()

cifar10_train = datasets.CIFAR10(data_path, train=True, download=dl,
↪transform=transform)
cifar10_test = datasets.CIFAR10(data_path, train=False, download=dl,
↪transform=transform)

batch_size = int(2**11)
workers = 12
cpu_prefetch = 39
gpu_prefetch = 28

print('begin init train_loader')
train_loader = DataLoader(
    cifar10_train,
    batch_size=batch_size,
    shuffle=True,
    num_workers=workers,
    prefetch_factor=cpu_prefetch,
    pin_memory=True
)

```



```

X_batch = next(iter(train_loader))[0]
dtype_size = X_batch.element_size()
print(f"Batch Size: {X_batch.element_size() * X_batch.nelement() / 1024**2}␣
↪MiB")

print('begin init fetcher')
train_loader_cuda = CudaDataPrefetcher(
    data_iterable = train_loader,
    device = torch.device('cuda'),
    num_prefetch_batches=gpu_prefetch
)
test_loader = DataLoader(cifar10_test, batch_size=len(cifar10_test),␣
↪shuffle=True, num_workers=workers, pin_memory=True, prefetch_factor=1)
test_loader_cuda = CudaDataPrefetcher(
    data_iterable = test_loader,
    device = torch.device('cuda'),
    num_prefetch_batches=1
)

model_1a = ConvImageClassifier(
    input_dim = 32,
    conv_layers=[32, 64],
    fc_layers=[32, 10],
    activation=nn.ReLU
).to(device=device)
params = sum(p.numel() for p in model_1a.parameters() if p.requires_grad)
print(f"Total parameters: {params}")
print(model_1a.stack)

model_1a.train_model(
    epochs=200,
    train_loader=train_loader_cuda,
    train_len=len(train_loader),
    test_loader=test_loader_cuda,
    test_len=len(test_loader),
    test_size = len(cifar10_test),
    loss_fn=nn.CrossEntropyLoss(),
    optimizer=torch.optim.Adam,
    optimizer_kwargs={'lr': 1e-3, 'weight_decay': 1e-2},
    print_epoch=1,
    header_epoch=15,
    #Disable scheduling
    sched_patience = 200
)

del train_loader

```

```
del test_loader
```

```
model_1a.plot_training("2 Layer CNN Training Curves")
```

```
begin init train_loader
```

```
Batch Size: 24.0 MiB
```

```
begin init fetcher
```

```
Total parameters: 150826
```

```
Sequential(
```

```
  (conv0): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
```

```
  (activation0): ReLU()
```

```
  (maxpool0): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,  
ceil_mode=False)
```

```
  (conv1): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
```

```
  (activation1): ReLU()
```

```
  (maxpool1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,  
ceil_mode=False)
```

```
  (flatten): Flatten(start_dim=1, end_dim=-1)
```

```
  (fc0): Linear(in_features=4096, out_features=32, bias=True)
```

```
  (activation_fc1): Tanh()
```

```
  (fc1): Linear(in_features=32, out_features=10, bias=True)
```

```
)
```

```
Training ConvImageClassifier
```

```
-----  
-----  
-----  
|      Epoch      | Epoch Time (s) | Training Loss | Test Loss  
| Overfit (%)     | Accuracy (%)   | Δ Accuracy (%) | GPU Memory  
(GiB) |
```

```
-----  
|      0          | 6.879051       | 1.937800      | 1.692992  
| -12.633269     | 42.219999     | 0.000000      | 4.213852  
|
```

```
-----  
|      1          | 4.994837       | 1.605037      | 1.515875  
| -5.555122      | 48.749998     | 15.466600     | 4.916979  
|
```

```
-----  
|      2          | 5.057521       | 1.465897      | 1.414413  
| -3.512099      | 51.569998     | 5.784616      | 5.618151  
|
```


	3		4.972616		1.375477		1.350838
	-1.791298		53.240001		3.238323		6.319323

	4		4.927670		1.308575		1.290943
	-1.347484		55.500001		4.244927		7.020495

	5		5.044797		1.260512		1.249116
	-0.904100		57.049996		2.792784		7.721667

	6		4.946557		1.217713		1.213817
	-0.319954		58.370000		2.313768		8.422839

	7		4.963850		1.191438		1.185038
	-0.537155		59.389997		1.747467		9.124011

	8		4.937151		1.141895		1.157622
	1.377297		59.560001		0.286251		9.825182

	9		4.935848		1.111418		1.121968
	0.949261		61.479998		3.223634		10.526354

	10		4.930698		1.078762		1.106483
	2.569699		62.089998		0.992192		11.227526

	11		4.939031		1.069479		1.099626
	2.818901		61.820000		-0.434848		11.928698
	12		4.886281		1.038594		1.061449
	2.200494		63.379997		2.077629		12.629870
	13		4.995658		1.013100		1.046757
	3.322247		63.709998		0.520671		13.331042
	14		4.999271		0.998499		1.042943
	4.451067		64.179999		0.737719		14.032214
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	15		4.955255		0.976865		1.017815
	4.191998		65.099996		1.433465		14.733386
	16		4.942564		0.963070		1.021136
	6.029279		64.489996		-0.937020		14.733386
	17		5.041818		0.945570		0.997559
	5.498090		66.139996		1.597542		14.733386

	18		5.018600		0.936561		0.982521
	4.907250		66.490000		0.529187		14.733386

	19		4.997144		0.920925		0.977600
	6.154187		66.719997		0.345911		14.733386

	20		5.019781		0.907129		0.972335
	7.188182		66.649997		-0.104916		14.733386

	21		5.009535		0.895712		0.975230
	8.877600		66.740000		0.029981		14.733386

	22		5.038126		0.891375		0.960982
	7.808997		67.089999		0.524421		14.733386

	23		5.222510		0.878327		0.942524
	7.308948		67.780000		1.028471		14.733386

	24		5.124631		0.869255		0.973694
	12.014804		66.719997		-1.563888		14.733386

	25		5.112703		0.867551		0.953081
	9.858801		67.009997		-1.136032		14.733386

	26		5.065426		0.851950		0.934716
	9.714802		68.000001		0.324581		14.733386
	27		5.058987		0.834209		0.945868
	13.385069		67.490000		-0.750002		14.733386
	28		5.017909		0.846333		0.925520
	9.356506		68.759996		1.117640		14.733386
	29		5.092650		0.823023		0.908128
	10.340587		69.000000		0.349046		14.733386
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	30		4.918151		0.822712		0.922056
	12.075120		68.349999		-0.942030		14.733386
	31		4.989425		0.818894		0.907556
	10.827121		69.330001		0.478262		14.733386
	32		4.983510		0.813875		0.901607
	10.779591		69.180000		-0.216358		14.733386

	33		5.009483		0.798466		0.910532
	14.035186		68.919998		-0.591378		14.733386
	34		4.931305		0.800569		0.899236
	12.324631		69.239998		-0.129818		14.733386
	35		5.006773		0.788863		0.893350
	13.245255		69.479996		0.216349		14.733386
	36		4.997746		0.796416		0.895868
	12.487496		69.229996		-0.359815		14.733386
	37		5.044682		0.796336		0.893554
	12.208151		69.309998		-0.244672		14.733386
	38		5.000170		0.779699		0.882630
	13.201449		69.849998		0.532530		14.733386
	39		4.996383		0.775605		0.881047
	13.594745		70.029998		0.257695		14.733386
	40		5.043076		0.776948		0.884165
	13.799740		69.940001		-0.128512		14.733386

	41		5.041635		0.773695		0.893078
	15.430255		69.239998		-1.128088		14.733386
	42		5.058732		0.770552		0.878690
	14.033852		70.269996		0.342708		14.733386
	43		5.043230		0.754261		0.870302
	15.384832		70.139998		-0.184997		14.733386
	44		4.998542		0.752734		0.874776
	16.213068		70.080000		-0.270380		14.733386
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	45		5.019987		0.743835		0.855158
	14.966059		71.209997		1.337699		14.733386
	46		5.004228		0.745706		0.872013
	16.937983		70.440000		-1.081304		14.733386
	47		5.011259		0.743847		0.855449
	15.003299		71.039999		-0.238728		14.733386

	48		4.993983		0.735898		0.864275
	17.444962		70.899999		-0.435329		14.733386
	49		5.076868		0.732867		0.872597
	19.066174		70.220000		-1.390251		14.733386
	50		5.019500		0.735458		0.866482
	17.815176		70.419997		-1.109395		14.733386
	51		4.996568		0.731122		0.863331
	18.082968		70.519996		-0.968967		14.733386
	52		4.973970		0.723913		0.852979
	17.829021		71.109998		-0.140428		14.733386
	53		4.986421		0.720687		0.850710
	18.041557		70.980000		-0.322983		14.733386
	54		5.057110		0.717377		0.881610
	22.893658		70.099998		-1.558769		14.733386
	55		5.026342		0.717588		0.856848
	19.406694		70.459998		-1.053222		14.733386

	56		5.025057		0.713417		0.882737
	23.733633		69.510001		-2.387300		14.733386
	57		5.034989		0.713721		0.839316
	17.597082		71.660000		0.631938		14.733386
	58		5.013689		0.704705		0.855010
	21.328808		70.669997		-1.381528		14.733386
	59		5.026370		0.708809		0.847124
	19.513639		71.270001		-0.544235		14.733386
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	60		4.995232		0.709052		0.869824
	22.674246		70.379996		-1.786218		14.733386
	61		5.050303		0.712813		0.855380
	20.000652		70.469999		-1.660621		14.733386
	62		5.017296		0.701495		0.868435
	23.797775		70.400000		-1.758304		14.733386

	63		5.017220		0.701392		0.841186
	19.931078		71.590000		-0.097683		14.733386
	64		5.014618		0.691178		0.828895
	19.924889		72.029996		0.516321		14.733386
	65		5.022207		0.691030		0.834924
	20.823246		71.590000		-0.610850		14.733386
	66		5.087819		0.679771		0.853366
	25.537222		71.079999		-1.318890		14.733386
	67		5.001268		0.692515		0.828224
	19.596493		72.139996		0.152715		14.733386
	68		4.985179		0.684203		0.820484
	19.918194		72.349995		0.291099		14.733386
	69		4.998172		0.680681		0.838111
	23.128300		71.329999		-1.409809		14.733386
	70		4.971375		0.687074		0.845096
	22.999332		71.399999		-1.313057		14.733386

	71		5.033550		0.677085		0.833596
	23.115574		71.980000		-0.511397		14.733386
	72		5.027353		0.674671		0.835942
	23.903578		71.669996		-0.939875		14.733386
	73		5.011705		0.674149		0.846878
	25.621727		71.109998		-1.713887		14.733386
	74		5.005163		0.669994		0.822554
	22.770489		72.249997		-0.138215		14.733386
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	75		5.051726		0.661650		0.827408
	25.052343		72.450000		0.138223		14.733386
	76		5.025922		0.671461		0.837118
	24.671137		71.599996		-1.173228		14.733386
	77		5.002988		0.672062		0.834897
	24.229179		71.619999		-1.145619		14.733386

	78		5.037694		0.674381		0.852034
	26.343037		71.029997		-1.959977		14.733386
	79		5.008406		0.673514		0.835284
	24.018772		71.689999		-1.049001		14.733386
	80		5.023103		0.664724		0.814567
	22.542067		72.329998		-0.165634		14.733386
	81		5.032135		0.660016		0.832222
	26.091066		71.520001		-1.283643		14.733386
	82		5.014982		0.658289		0.827561
	25.713850		72.069997		-0.524504		14.733386
	83		4.990252		0.657215		0.830514
	26.368765		72.169995		-0.386480		14.733386
	84		5.049405		0.646537		0.816893
	26.349035		72.529995		0.110415		14.733386
	85		5.018591		0.647542		0.815600
	25.953310		72.700000		0.234392		14.733386

	86		5.033751		0.647618		0.812683
	25.488043		72.389996		-0.426415		14.733386
	87		5.045481		0.646377		0.815934
	26.231840		72.779995		0.110035		14.733386
	88		4.996143		0.648928		0.810473
	24.894083		72.670001		-0.151133		14.733386
	89		5.029956		0.639077		0.808710
	26.543398		73.259997		0.659525		14.733386
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	90		5.110249		0.641995		0.831552
	29.526297		71.480000		-2.429699		14.733386
	91		4.982387		0.649811		0.825403
	27.022100		71.929997		-1.815452		14.733386
	92		4.985412		0.635004		0.812187
	27.902637		72.479999		-1.064699		14.733386

	93		5.036469		0.644235		0.817088
	26.830687		72.169995		-1.487854		14.733386
	94		5.012946		0.632602		0.819317
	29.515427		71.819997		-1.965603		14.733386
	95		5.045994		0.633972		0.844893
	33.269829		70.599997		-3.630904		14.733386
	96		4.988159		0.644581		0.841811
	30.598265		71.309996		-2.661755		14.733386
	97		5.099520		0.642988		0.819964
	27.524033		72.079998		-1.610700		14.733386
	98		5.076290		0.630325		0.864446
	37.142879		70.159996		-4.231507		14.733386
	99		5.029150		0.642412		0.865549
	34.734249		69.690001		-4.873050		14.733386
	100		5.074078		0.639341		0.831916
	30.120871		71.419996		-2.511605		14.733386

	101		5.076982		0.641404		0.846540				
	31.982470		70.739996		-3.439805		14.733386				
	102		5.054729		0.635157		0.799278				
	25.839453		72.939998		-0.436800		14.733386				
	103		5.141601		0.625099		0.819944				
	31.170256		72.029996		-1.678954		14.733386				
	104		5.070658		0.622841		0.818336				
	31.387720		72.420001		-1.146596		14.733386				
	Epoch		Epoch Time (s)		Training Loss		Test Loss				
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory				
(GiB)											
	105		5.019599		0.628627		0.813345				
	29.384373		72.560000		-0.955497		14.733386				
	106		5.024177		0.626152		0.827057				
	32.085606		71.630001		-2.224947		14.733386				
	107		5.013819		0.629211		0.829194				
	31.783247		71.639997		-2.211303		14.733386				

	108		4.982967		0.627714		0.834247
	32.902443		71.789998		-2.006552		14.733386
	109		5.065071		0.623302		0.814757
	30.716171		72.819996		-0.600603		14.733386
	110		5.052822		0.620189		0.816181
	31.602104		72.520000		-1.010098		14.733386
	111		5.039487		0.620462		0.830820
	33.903545		71.509999		-2.388750		14.733386
	112		5.014404		0.628477		0.807500
	28.485163		72.859997		-0.546002		14.733386
	113		5.031634		0.607373		0.806208
	32.736797		72.799999		-0.627899		14.733386
	114		5.058440		0.613070		0.832420
	35.778915		71.499997		-2.402402		14.733386
	115		5.137480		0.615059		0.818509
	33.077972		72.349995		-1.242154		14.733386

	116		5.018563		0.609816		0.812268
	33.198799		72.399998		-1.173901		14.733386
	117		5.109816		0.607323		0.806634
	32.818096		72.839999		-0.573298		14.733386
	118		5.038103		0.607644		0.810839
	33.439774		72.499996		-1.037402		14.733386
	119		5.046576		0.606106		0.822247
	35.660679		71.880001		-1.883697		14.733386
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	120		5.127125		0.604725		0.795232
	31.503040		73.170000		-0.122846		14.733386
	121		5.092554		0.603490		0.803289
	33.107285		72.639996		-0.846303		14.733386
	122		5.157858		0.607773		0.808804
	33.076607		72.589999		-0.914548		14.733386

	123		5.063401		0.606742		0.810017
	33.502678		72.380000		-1.201197		14.733386
	124		5.077441		0.600606		0.790615
	31.636356		73.249996		-0.013652		14.733386
	125		5.053849		0.614245		0.812236
	32.233276		72.359997		-1.228502		14.733386
	126		5.084872		0.599685		0.803532
	33.992432		72.589999		-0.914548		14.733386
	127		5.167486		0.599114		0.815815
	36.170300		72.490001		-1.051046		14.733386
	128		5.089398		0.602681		0.808109
	34.085697		72.829998		-0.586950		14.733386
	129		5.076684		0.600694		0.800499
	33.262379		73.109996		-0.204752		14.733386
	130		5.027541		0.597062		0.817026
	36.840973		71.969998		-1.760851		14.733386

	131		5.041021		0.594622		0.802022
	34.879238		72.869998		-0.532349		14.733386
	132		5.119192		0.598600		0.819156
	36.845387		71.880001		-1.883697		14.733386
	133		5.096864		0.595433		0.801509
	34.609364		73.139995		-0.163803		14.733386
	134		5.065614		0.587363		0.809745
	37.861038		72.679996		-0.791702		14.733386
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	135		5.141582		0.584925		0.789208
	34.924698		73.619998		0.491401		14.733386
	136		5.089416		0.588752		0.831994
	41.314877		71.779996		-2.499323		14.733386
	137		5.073261		0.602229		0.816510
	35.581394		72.209996		-1.915243		14.733386

	138		5.051284		0.588933		0.806700
	36.976452		72.749996		-1.181746		14.733386
	139		5.172863		0.596244		0.793793
	33.132214		73.219997		-0.543332		14.733386
	140		5.030095		0.589227		0.799713
	35.722385		72.749996		-1.181746		14.733386
	141		5.064724		0.587496		0.805443
	37.097630		72.929996		-0.937247		14.733386
	142		5.317110		0.589445		0.804893
	36.550980		72.659999		-1.303991		14.733386
	143		5.075824		0.595370		0.822761
	38.193150		72.189999		-1.942405		14.733386
	144		5.127365		0.595650		0.820195
	37.697517		72.229999		-1.888072		14.733386
	145		5.070822		0.589687		0.787928
	33.618114		73.519999		-0.135831		14.733386

	146		5.024504		0.584502		0.831057
	42.182186		71.639997		-2.689488		14.733386
	147		5.040956		0.586660		0.796234
	35.723175		72.999996		-0.842165		14.733386
	148		5.033256		0.580001		0.790484
	36.290077		73.579997		-0.054334		14.733386
	149		5.084164		0.583147		0.831220
	42.540440		71.969998		-2.241238		14.733386
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	150		5.084607		0.597433		0.805152
	34.768665		72.340000		-1.738655		14.733386
	151		5.077116		0.578795		0.797499
	37.786106		72.889996		-0.991581		14.733386
	152		5.062029		0.580120		0.799035
	37.736263		73.049998		-0.774245		14.733386

	153		5.031142		0.578993		0.802004
	38.517136		72.829998		-1.073078		14.733386

	154		5.084552		0.580315		0.816360
	40.675404		72.499996		-1.521327		14.733386

	155		5.067575		0.578752		0.804299
	38.971237		72.380000		-1.684321		14.733386

	156		5.113379		0.579363		0.794566
	37.144699		73.429996		-0.258084		14.733386

	157		5.017329		0.578061		0.798019
	38.051037		73.069996		-0.747082		14.733386

	158		5.101545		0.575089		0.804189
	39.837406		72.529995		-1.480579		14.733386

	159		5.028152		0.571383		0.781985
	36.858311		73.569995		-0.067920		14.733386

	160		5.082304		0.580265		0.817013
	40.800018		72.259998		-1.847323		14.733386

	161		5.036721		0.578746		0.811182
	40.162014		72.319996		-1.765826		14.733386
	162		5.045552		0.588903		0.808974
	37.369617		72.439998		-1.602824		14.733386
	163		5.086725		0.584779		0.819144
	40.077641		72.369999		-1.697907		14.733386
	164		5.100532		0.570557		0.798689
	39.984184		72.950000		-0.910076		14.733386
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	165		5.097234		0.566995		0.822130
	44.997696		72.149998		-1.996740		14.733386
	166		5.046937		0.571805		0.786729
	37.586933		73.299998		-0.434664		14.733386
	167		5.059964		0.572341		0.806426
	40.899544		72.649997		-1.317577		14.733386

	168		5.083268		0.578109		0.803447
	38.978401		72.819996		-1.086663		14.733386
	169		5.034514		0.575779		0.808265
	40.377739		72.499996		-1.521327		14.733386
	170		5.096079		0.571691		0.840251
	46.976452		71.059996		-3.477318		14.733386
	171		5.090700		0.577162		0.804536
	39.395245		72.749996		-1.181746		14.733386
	172		5.041945		0.570497		0.800041
	40.235794		72.889996		-0.991581		14.733386
	173		5.052004		0.574199		0.788621
	37.342800		73.259997		-0.488998		14.733386
	174		5.052775		0.572699		0.799113
	39.534668		73.420000		-0.271662		14.733386
	175		5.074750		0.567056		0.807956
	42.482597		72.109997		-2.051074		14.733386

	176		5.063622		0.566776		0.795557
	40.365253		73.119998		-0.679163		14.733386
	177		5.112973		0.560236		0.817368
	45.897247		71.889997		-2.349906		14.733386
	178		5.141950		0.579844		0.846047
	45.909477		71.459997		-2.933986		14.733386
	179		5.097329		0.579943		0.794550
	37.004822		73.069996		-0.747082		14.733386
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	180		5.048506		0.560042		0.789484
	40.968742		73.219997		-0.543332		14.733386
	181		5.058635		0.556483		0.799255
	43.626152		72.679996		-1.276828		14.733386
	182		5.078006		0.561945		0.807073
	43.621426		72.779995		-1.140997		14.733386


	183		5.123325		0.559224		0.810336
	44.903713		72.469997		-1.562076		14.733386
	184		5.105889		0.562512		0.792792
	40.937675		73.049998		-0.774245		14.733386
	185		5.135925		0.562172		0.794817
	41.383297		73.170000		-0.611243		14.733386
	186		5.117503		0.562027		0.808710
	43.891663		72.289997		-1.806575		14.733386
	187		5.105713		0.560768		0.796730
	42.078255		72.819996		-1.086663		14.733386
	188		5.096840		0.558971		0.793019
	41.871277		72.889996		-0.991581		14.733386
	189		5.079002		0.567355		0.829393
	46.185753		71.509999		-2.866067		14.733386
	190		5.047444		0.560887		0.794382
	41.629650		72.979999		-0.869328		14.733386

	191		5.137882		0.566190		0.788644
	39.289711		73.119998		-0.679163		14.733386
	192		5.176906		0.564188		0.819275
	45.213093		72.579998		-1.412659		14.733386
	193		5.095566		0.558991		0.790708
	41.452759		73.390001		-0.312410		14.733386
	194		5.114154		0.560496		0.800392
	42.800682		72.979999		-0.869328		14.733386
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	195		5.104700		0.553329		0.793020
	43.317894		72.819996		-1.086663		14.733386
	196		5.060708		0.560255		0.796791
	42.219360		72.920001		-0.950824		14.733386
	197		5.068001		0.556775		0.787571
	41.452240		73.390001		-0.312410		14.733386

	198		5.114279		0.551170		0.818402
	48.484344		72.380000		-1.684321		14.733386

	199		5.161066		0.557756		0.828847
	48.603931		71.619999		-2.716650		14.733386

Training Time: 465.7891447544098 seconds



```

[7]: try:
    del train_loader
    del test_loader
    del model_1a
    del model_1b
    del resnet
    del train_loader_cuda
    del test_loader_cuda
except:
    pass

# Reset CUDA context
start = time.time()
torch.cuda.empty_cache()
torch.cuda.reset_peak_memory_stats()

gc.collect()

cifar10_train = datasets.CIFAR10(data_path, train=True, download=dl,
    ↪transform=transform)
cifar10_test = datasets.CIFAR10(data_path, train=False, download=dl,
    ↪transform=transform)

batch_size = int(2**11)
workers = 12
cpu_prefetch = 39
gpu_prefetch = 28

print('begin init train_loader')
train_loader = DataLoader(
    cifar10_train,
    batch_size=batch_size,
    shuffle=True,
    num_workers=workers,
    prefetch_factor=cpu_prefetch,
    pin_memory=True
)

X_batch = next(iter(train_loader))[0]
dtype_size = X_batch.element_size()
print(f"Batch Size: {X_batch.element_size() * X_batch.nelement() / 1024**2}
    ↪MiB")

print('begin init fetcher')
train_loader_cuda = CudaDataPrefetcher(

```

```

        data_iterable = train_loader,
        device = torch.device('cuda'),
        num_prefetch_batches=gpu_prefetch
    )
test_loader = DataLoader(cifar10_test, batch_size=len(cifar10_test),
    ↪shuffle=True, num_workers=workers, pin_memory=True, prefetch_factor=1)
test_loader_cuda = CudaDataPrefetcher(
    data_iterable = test_loader,
    device = torch.device('cuda'),
    num_prefetch_batches=1
)

model_1b = ConvImageClassifier(
    input_dim = 32,
    conv_layers=[32, 64, 128],
    fc_layers=[32, 10],
    activation=nn.ReLU
).to(device=device)
params = sum(p.numel() for p in model_1b.parameters() if p.requires_grad)
print(f"Total parameters: {params}")

print(model_1b.stack)

model_1b.train_model(
    epochs=200,
    train_loader=train_loader_cuda,
    train_len=len(train_loader),
    test_loader=test_loader_cuda,
    test_len=len(test_loader),
    test_size = len(cifar10_test),
    optimizer = torch.optim.Adam,
    optimizer_kwargs={'lr': 4e-4, 'weight_decay': 1e-2}, #Increase alpha to 2
    ↪next time
    loss_fn=nn.CrossEntropyLoss(),
    print_epoch=1,
    #Disable scheduling
    sched_patience = 200
)
del train_loader
del test_loader
model_1b.plot_training("3 Layer CNN Training Curves")

```

```

begin init train_loader
Batch Size: 24.0 MiB
begin init fetcher
Sequential(
  (conv0): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

```

```

(activation0): ReLU()
(maxpool0): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(conv1): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(activation1): ReLU()
(maxpool1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(conv2): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(activation2): ReLU()
(maxpool2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(flatten): Flatten(start_dim=1, end_dim=-1)
(fc0): Linear(in_features=2048, out_features=32, bias=True)
(activation_fc1): Tanh()
(fc1): Linear(in_features=32, out_features=10, bias=True)
)

```

Training ConvImageClassifier

Epoch	Epoch Time (s)	Training Loss	Test Loss
Overfit (%)	Accuracy (%)	Δ Accuracy (%)	GPU Memory (GiB)
0	6.599083	2.123456	1.948487
-8.239858	30.339998	0.000000	4.397570
1	5.095938	1.884181	1.808225
-4.031254	36.960000	21.819387	5.098743
2	5.104386	1.763717	1.702580
-3.466377	41.149998	11.336574	5.799916
3	5.259627	1.671804	1.621082
-3.033968	43.860000	6.585669	6.501088

	4		5.091378		1.603726		1.563755
	-2.492332		45.389998		3.488367		7.202260

	5		5.196554		1.546524		1.518240
	-1.828888		47.219998		4.031724		7.903431

	6		5.055001		1.505353		1.491611
	-0.912913		47.829998		1.291825		8.604603

	7		5.300217		1.466276		1.470768
	0.306398		47.700000		-0.271791		9.305775

	8		5.158196		1.436028		1.435479
	-0.038194		50.019997		4.578716		10.006947

	9		5.092443		1.406422		1.398012
	-0.598020		51.370001		2.698928		10.708119

	10		5.108890		1.377094		1.370624
	-0.469776		52.129996		1.479453		11.409291

	11		5.151833		1.352459		1.352640
	0.013336		52.759999		1.208523		12.110463

	12		5.107094		1.329279		1.321973
	-0.549647		53.729999		1.838514		12.811635
	13		5.113696		1.310188		1.304062
	-0.467534		54.240000		0.949193		13.512806
	14		5.095271		1.286842		1.288354
	0.117547		55.339998		2.028019		14.213978
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	15		5.078908		1.272123		1.276889
	0.374658		55.229998		-0.198772		14.915150
	16		5.201680		1.258695		1.267410
	0.692424		56.049997		1.282975		14.915150
	17		5.199535		1.244539		1.255717
	0.898146		56.430000		0.677972		14.915150
	18		5.049537		1.226315		1.230540
	0.344559		57.080001		1.151871		14.915150

	19		5.060830		1.211456		1.219420
	0.657372		57.639998		0.981075		14.915150
	20		5.101258		1.194798		1.198727
	0.328824		58.309996		1.162384		14.915150
	21		5.119519		1.180924		1.191759
	0.917526		58.340001		0.051458		14.915150
	22		5.043655		1.173767		1.185677
	1.014699		58.649999		0.531364		14.915150
	23		5.075843		1.161882		1.175911
	1.207450		58.849996		0.341002		14.915150
	24		5.146863		1.150066		1.162871
	1.113444		59.700000		1.444357		14.915150
	25		5.158943		1.138020		1.153674
	1.375481		59.990001		0.485763		14.915150
	26		5.087031		1.123744		1.139605
	1.411456		60.240000		0.416736		14.915150

	27		5.090375		1.116595		1.151081
	3.088480		59.689999		-0.913018		14.915150
	28		5.113190		1.110461		1.126441
	1.439053		60.899997		1.095611		14.915150
	29		5.114990		1.096704		1.135921
	3.575929		60.310000		-0.968796		14.915150
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	30		5.186446		1.088094		1.108954
	1.917087		61.430001		0.870287		14.915150
	31		5.093885		1.077307		1.121839
	4.133659		60.749996		-1.106960		14.915150
	32		5.126390		1.076997		1.099699
	2.107863		62.000000		0.927884		14.915150
	33		5.127231		1.064820		1.102011
	3.492626		61.299998		-1.129037		14.915150

	34		5.073361		1.051656		1.089044
	3.555117		62.140000		0.225805		14.915150

	35		5.091969		1.046859		1.067175
	1.940709		63.190001		1.689735		14.915150

	36		5.078733		1.033740		1.061945
	2.728457		63.459998		0.427278		14.915150

	37		5.084619		1.026881		1.061484
	3.369730		63.519996		0.094545		14.915150

	38		5.107906		1.016998		1.053293
	3.568783		63.290000		-0.362085		14.915150

	39		5.128337		1.009513		1.045956
	3.609933		63.589996		0.110201		14.915150

	40		5.097540		0.998343		1.029674
	3.138276		64.639997		1.651205		14.915150

	41		5.108642		0.990487		1.027441
	3.730927		64.609998		-0.046409		14.915150

	42		5.131843		0.980696		1.030523
	5.080703		64.289999		-0.541458		14.915150
	43		5.013742		0.982274		1.052352
	7.134252		63.379997		-1.949258		14.915150
	44		5.081797		0.978605		1.017990
	4.024603		65.050000		0.634286		14.915150
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	45		5.021367		0.966912		1.030294
	6.555111		64.709997		-0.522679		14.915150
	46		5.044518		0.958193		1.009344
	5.338353		65.270001		0.338203		14.915150
	47		5.089929		0.951124		1.005510
	5.718028		65.130001		-0.214493		14.915150
	48		5.070228		0.944371		0.995486
	5.412549		65.899998		0.965217		14.915150

	49		5.104623		0.945438		1.011090
	6.944112		64.840001		-1.608493		14.915150

	50		5.083852		0.937149		0.984229
	5.023837		66.219997		0.485583		14.915150

	51		5.122586		0.925459		0.984319
	6.360085		66.249996		0.045302		14.915150

	52		5.096201		0.922812		0.978621
	6.047788		65.990001		-0.392446		14.915150

	53		5.131413		0.921411		0.997682
	8.277654		65.520000		-1.101881		14.915150

	54		5.056648		0.908973		0.960075
	5.621919		67.030001		1.177365		14.915150

	55		5.029018		0.903034		0.963079
	6.649313		67.140001		0.164106		14.915150

	56		5.040206		0.902585		0.951875
	5.460921		67.420000		0.417037		14.915150

	57		5.028303		0.893709		0.948499
	6.130551		67.170000		-0.370809		14.915150
	58		5.011108		0.891331		0.965768
	8.351196		66.469997		-1.409081		14.915150
	59		5.085455		0.889393		0.949489
	6.756923		67.379999		-0.059331		14.915150
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	60		5.060808		0.885069		0.935363
	5.682459		67.750001		0.489471		14.915150
	61		5.040805		0.877112		0.930217
	6.054521		68.099999		0.516603		14.915150
	62		5.088319		0.869066		0.938141
	7.948235		67.890000		-0.308369		14.915150
	63		5.085271		0.873677		0.937699
	7.327950		67.809999		-0.425845		14.915150

	64		5.123614		0.868866		0.923938
	6.338306		68.500000		0.587373		14.915150
	65		5.134030		0.859185		0.926281
	7.809257		68.110001		-0.569342		14.915150
	66		5.078809		0.857151		0.925934
	8.024608		68.129998		-0.540149		14.915150
	67		5.098463		0.854254		0.924849
	8.264010		68.229997		-0.394165		14.915150
	68		5.111977		0.859365		0.935850
	8.900144		68.629998		0.189778		14.915150
	69		5.039732		0.840639		0.929612
	10.584030		68.009996		-0.903397		14.915150
	70		5.085504		0.838153		0.907637
	8.290179		69.000000		0.539125		14.915150
	71		5.067537		0.843908		0.929506
	10.143097		67.820001		-1.710144		14.915150

	72		5.057573		0.835999		0.915937
	9.561942		68.529999		-0.681160		14.915150
	73		5.074442		0.837075		0.904336
	8.035266		69.199997		0.289851		14.915150
	74		5.095749		0.832735		0.917949
	10.233088		68.949997		-0.361271		14.915150
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	75		5.059535		0.821980		0.893973
	8.758437		69.690001		0.708098		14.915150
	76		5.055362		0.817010		0.903975
	10.644279		69.389999		-0.430481		14.915150
	77		5.104808		0.819375		0.889156
	8.516319		69.499999		-0.272638		14.915150
	78		5.068392		0.813133		0.900651
	10.763136		69.229996		-0.660073		14.915150

	79		5.061541		0.808862		0.908639
	12.335466		68.930000		-1.090545		14.915150
	80		5.071167		0.806705		0.894210
	10.847201		69.510001		-0.258287		14.915150
	81		5.105503		0.804085		0.888455
	10.492674		69.369996		-0.459184		14.915150
	82		5.102875		0.801373		0.897731
	12.024125		69.349998		-0.487879		14.915150
	83		5.091874		0.801492		0.882098
	10.056890		70.269996		0.832250		14.915150
	84		5.102546		0.796571		0.907992
	13.987577		68.909997		-1.935391		14.915150
	85		5.091497		0.793067		0.872336
	9.995210		70.429999		0.227697		14.915150
	86		5.108479		0.786681		0.880855
	11.971060		69.870001		-0.795112		14.915150

	87		5.188274		0.782331		0.870107
	11.219831		70.489997		0.085188		14.915150
	88		5.092411		0.778280		0.871861
	12.024100		70.300001		-0.269536		14.915150
	89		5.157857		0.779800		0.886968
	13.743012		70.120001		-0.524891		14.915150
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	90		5.120567		0.777302		0.886984
	14.110626		69.830000		-0.936298		14.915150
	91		5.096100		0.774991		0.875539
	12.974070		70.400000		-0.127674		14.915150
	92		5.158041		0.780365		0.866458
	11.032430		70.469999		-0.028369		14.915150
	93		5.121035		0.781344		0.915482
	17.167654		68.439996		-2.908215		14.915150

	94		5.098719		0.782252		0.879924
	12.485970		70.099998		-0.553269		14.915150

	95		5.121605		0.769990		0.859838
	11.668668		70.739996		0.354660		14.915150

	96		5.093006		0.763771		0.871130
	14.056423		70.559996		-0.254453		14.915150

	97		5.114752		0.765917		0.874804
	14.216483		69.999999		-1.046081		14.915150

	98		5.081268		0.765277		0.870959
	13.809715		70.349997		-0.551314		14.915150

	99		5.104732		0.758171		0.858317
	13.208964		70.539999		-0.282722		14.915150

	100		5.126863		0.755749		0.852160
	12.756919		71.359998		0.876451		14.915150

	101		5.160976		0.756665		0.870397
	15.030764		70.319998		-1.457399		14.915150

	102		5.071728		0.754289		0.870921
	15.462502		70.480001		-1.233180		14.915150
	103		5.068205		0.752413		0.856732
	13.864575		70.940000		-0.588562		14.915150
	104		5.067548		0.750012		0.849197
	13.224454		71.039999		-0.448430		14.915150
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	105		5.069998		0.740821		0.860605
	16.169134		70.489997		-1.219172		14.915150
	106		5.074821		0.751605		0.856403
	13.943147		71.129996		-0.322313		14.915150
	107		5.041211		0.748542		0.849530
	13.491354		71.069998		-0.406391		14.915150
	108		5.067728		0.735770		0.846414
	15.037781		71.219999		-0.196187		14.915150

	109		5.120624		0.735236		0.834490
	13.499506		71.359998		0.000000		14.915150

	110		5.120245		0.736660		0.850646
	15.473345		71.099997		-0.364352		14.915150

	111		5.119845		0.743295		0.881617
	18.609299		69.569999		-2.508407		14.915150

	112		5.103891		0.732732		0.839506
	14.571980		71.499997		0.196187		14.915150

	113		5.176232		0.722651		0.838164
	15.984736		71.480000		-0.027968		14.915150

	114		5.130633		0.720923		0.845920
	17.338436		71.389997		-0.153847		14.915150

	115		5.117654		0.726849		0.847043
	16.536243		70.980000		-0.727268		14.915150

	116		5.059834		0.722787		0.836406
	15.719629		71.700001		0.279725		14.915150

	117		5.077767		0.716625		0.839947
	17.208736		71.789998		0.125519		14.915150
	118		5.042213		0.711941		0.829023
	16.445593		72.219998		0.598969		14.915150
	119		5.109068		0.710683		0.835860
	17.613676		71.770000		-0.623092		14.915150
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	120		5.136141		0.708005		0.823399
	16.298500		72.139996		-0.110775		14.915150
	121		5.081880		0.707609		0.822545
	16.242987		72.119999		-0.138464		14.915150
	122		5.100162		0.703986		0.853264
	21.204630		70.879996		-1.855444		14.915150
	123		5.043020		0.721448		0.820083
	13.671699		72.099996		-0.166162		14.915150

	124		5.151543		0.705544		0.819741
	16.185585		72.329998		0.152313		14.915150
	125		5.112077		0.697202		0.821662
	17.851221		72.279996		-0.069131		14.915150
	126		5.074937		0.706385		0.828553
	17.294765		71.759999		-0.788054		14.915150
	127		5.112021		0.712020		0.842574
	18.335752		71.829998		-0.691275		14.915150
	128		5.068683		0.701907		0.842251
	19.994738		71.429998		-1.244298		14.915150
	129		5.106611		0.703883		0.825850
	17.327799		71.730000		-0.829529		14.915150
	130		5.161842		0.693148		0.821358
	18.496904		72.149998		-0.248859		14.915150
	131		5.114609		0.692579		0.835119
	20.580921		71.469998		-1.188994		14.915150

	132		5.234236		0.694519		0.826709
	19.033396		72.380000		0.069131		14.915150
	133		5.167754		0.683442		0.818009
	19.689714		72.240001		-0.193423		14.915150
	134		5.169732		0.686514		0.827337
	20.512833		72.119999		-0.359217		14.915150
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	135		5.170857		0.685092		0.819798
	19.662561		72.240001		-0.193423		14.915150
	136		5.165364		0.681307		0.813139
	19.349766		72.450000		0.096711		14.915150
	137		5.184544		0.676408		0.827789
	22.380207		71.980000		-0.648724		14.915150
	138		5.041432		0.682127		0.816983
	19.769909		72.209996		-0.331269		14.915150

	139		5.188199		0.672820		0.811393
	20.595865		72.950000		0.690130		14.915150
	140		5.134880		0.672488		0.810476
	20.519125		72.719997		-0.315288		14.915150
	141		5.105504		0.675568		0.805462
	19.227417		72.899997		-0.068543		14.915150
	142		5.120533		0.671230		0.808244
	20.412378		72.969997		0.027412		14.915150
	143		5.068957		0.665354		0.813597
	22.280344		72.499996		-0.644101		14.915150
	144		5.135123		0.667591		0.809880
	21.313772		72.920001		-0.068516		14.915150
	145		5.082710		0.663192		0.802908
	21.067223		72.700000		-0.370011		14.915150
	146		5.110504		0.666857		0.804756
	20.678925		72.799999		-0.232970		14.915150

	147		5.176138		0.662117		0.819327
	23.743467		72.689998		-0.383718		14.915150
	148		5.138168		0.666574		0.826696
	24.021763		71.569997		-1.918597		14.915150
	149		5.160818		0.664240		0.801299
	20.634020		73.249996		0.383718		14.915150
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	150		5.115169		0.659333		0.801870
	21.618420		73.039997		-0.286688		14.915150
	151		5.125012		0.653944		0.823548
	25.935585		72.270000		-1.337878		14.915150
	152		5.116081		0.655704		0.795027
	21.247782		73.100001		-0.204771		14.915150
	153		5.120463		0.658052		0.818986
	24.456205		72.240001		-1.378832		14.915150

	154		5.186667		0.649822		0.794072
	22.198418		73.359996		0.150171		14.915150

	155		5.152768		0.646464		0.794968
	22.971737		73.039997		-0.436204		14.915150

	156		5.118167		0.651932		0.799742
	22.672680		73.259997		-0.136312		14.915150

	157		5.073302		0.654624		0.801499
	22.436533		72.939998		-0.572517		14.915150

	158		5.136588		0.644021		0.797534
	23.836645		73.219997		-0.190839		14.915150

	159		5.128449		0.647507		0.798428
	23.307957		73.100001		-0.354410		14.915150

	160		5.112539		0.641143		0.792088
	23.543064		73.240000		-0.163572		14.915150

	161		5.120852		0.643312		0.809783
	25.877104		72.719997		-0.872409		14.915150

	162		5.094538		0.641172		0.794955
	23.984808		73.299998		-0.081786		14.915150
	163		5.057532		0.635318		0.784141
	23.424995		73.679996		0.436204		14.915150
	164		5.093210		0.631926		0.785027
	24.227589		73.759997		0.108579		14.915150
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	165		5.111727		0.631431		0.796851
	26.197598		73.299998		-0.623643		14.915150
	166		5.123473		0.634222		0.789436
	24.473181		73.460001		-0.406719		14.915150
	167		5.107878		0.634861		0.785529
	23.732395		73.579997		-0.244035		14.915150
	168		5.150738		0.629582		0.806329
	28.073561		72.909999		-1.152383		14.915150

	169		5.091628		0.628832		0.790047
	25.637213		73.670000		-0.122013		14.915150
	170		5.073217		0.635245		0.781749
	23.062716		73.909998		0.203364		14.915150
	171		5.121365		0.632413		0.813234
	28.592190		72.889996		-1.380060		14.915150
	172		5.029716		0.631198		0.793187
	25.663792		73.299998		-0.825328		14.915150
	173		5.140740		0.633403		0.781718
	23.415546		73.589998		-0.432958		14.915150
	174		5.088692		0.627700		0.786654
	25.323292		73.569995		-0.460023		14.915150
	175		5.156265		0.625505		0.783891
	25.321344		73.490000		-0.568256		14.915150
	176		5.202651		0.622379		0.784202
	26.000751		73.490000		-0.568256		14.915150


	177		5.100004		0.622683		0.793957
	27.505703		73.149997		-1.028279		14.915150
	178		5.156939		0.618830		0.775399
	25.300802		74.180001		0.365313		14.915150
	179		5.132620		0.615692		0.781508
	26.931622		73.839998		-0.458348		14.915150
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	180		5.125704		0.614744		0.775486
	26.147738		73.939997		-0.323543		14.915150
	181		5.205902		0.616690		0.781392
	26.707546		73.859996		-0.431390		14.915150
	182		5.122013		0.629987		0.779209
	23.686480		73.759997		-0.566196		14.915150
	183		5.205942		0.611550		0.774581
	26.658596		73.829997		-0.471831		14.915150

	184		5.180202		0.608123		0.788122
	29.599031		73.390001		-1.064977		14.915150
	185		5.212535		0.609963		0.781101
	28.057093		73.789996		-0.525755		14.915150
	186		5.164953		0.605180		0.816878
	34.980904		72.209996		-2.655709		14.915150
	187		5.107094		0.615284		0.779685
	26.719515		73.890001		-0.390942		14.915150
	188		5.131805		0.610408		0.776326
	27.181515		73.939997		-0.323543		14.915150
	189		5.131907		0.609715		0.777072
	27.448435		74.009997		-0.229178		14.915150
	190		5.079273		0.605033		0.787273
	30.120705		73.699999		-0.647078		14.915150
	191		5.156400		0.601853		0.791279
	31.473764		72.920001		-1.698572		14.915150

	192		5.183256		0.597795		0.768771
	28.601042		74.379998		0.269611		14.915150
	193		5.132535		0.591397		0.770208
	30.235384		74.210000		-0.228554		14.915150
	194		5.130533		0.598124		0.777528
	29.994465		73.850000		-0.712555		14.915150
	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							
	195		5.087523		0.592661		0.776363
	30.996077		74.100000		-0.376444		14.915150
	196		5.129474		0.594608		0.786006
	32.188976		73.569995		-1.089007		14.915150
	197		5.143143		0.599006		0.777057
	29.724390		73.920000		-0.618444		14.915150
	198		5.106315		0.599403		0.778661
	29.906042		73.749995		-0.847006		14.915150

	199		5.109549		0.591137		0.773772
	30.895559		74.309999		-0.094111		14.915150

Training Time: 472.81531620025635 seconds



[4]: *#after last bn but before last weight*

```
class ResBlock(nn.Module):
    def __init__(self, in_chans, out_chans, nonlinearity = 'relu', stride=1,
        ↪ dropout = 0.4):
        super().__init__()
```

```

        self.conv1 = nn.Conv2d(in_chans, out_chans, kernel_size=3, padding=1,
↪bias=False, stride=stride)
        self.batch_norm1 = nn.BatchNorm2d(num_features=out_chans)
        self.conv2 = nn.Conv2d(out_chans, out_chans, kernel_size=3, padding=1,
↪bias=False, stride=stride)
        self.batch_norm2 = nn.BatchNorm2d(num_features=out_chans)
        self.dropout = nn.Dropout(dropout)
        self.shortcut = nn.Conv2d(in_chans, out_chans, kernel_size=1, stride=1,
↪bias=False) if in_chans != out_chans else nn.Identity()

        torch.nn.init.kaiming_normal_(self.conv1.weight,
↪nonlinearity=nonlinearity)
        torch.nn.init.constant_(self.batch_norm1.weight, 0.5)
        torch.nn.init.zeros_(self.batch_norm1.bias)
        torch.nn.init.kaiming_normal_(self.conv2.weight,
↪nonlinearity=nonlinearity)
        torch.nn.init.zeros_(self.batch_norm2.bias)
    def forward(self, x):
        out = self.batch_norm1(x)
        out = F.relu(out)
        out = self.conv1(out)
        out = self.batch_norm2(out)
        out = F.relu(out)
        out = self.dropout(out)
        out = self.conv2(out)
        out += self.shortcut(x)
        return out
class ResNet(Classifier):
    def __init__(self, input_dim = 32, n_blocks = 10, conv_channels = [32,16],
↪fc_channels = [32, 10], dropout_p=0.4, dropout_h=0.4, nonlinearity='relu'):
        super().__init__()

        # Add initial convolutions
        self.h1 = nn.Sequential()
        for i in range(len(conv_channels)):
            self.h1.add_module(
                name=f'conv{i}',
                module=nn.Conv2d(
                    in_channels = 3 if i == 0 else conv_channels[i-1],
                    out_channels=conv_channels[i],
                    kernel_size=3,
                    padding=1
                )
            )
        self.h1.add_module(

```

```

        name=f'activation{i}',
        module=nn.ReLU()
    )
    self.h1.add_module(
        name=f'maxpool',
        module=nn.MaxPool2d(2)
    )
    #output of h1 before maxpool is 16 32x32 images. After maxpool, 16
    ↪ 16x16 images
    # h1_in: torch.Size([1024, 3, 32, 32])
    # res_block_in: torch.Size([1024, 16, 16, 16])
    # h2_in: torch.Size([1024, 64])
    #Add Resblocks
    res_block_in = conv_channels[0]//2
    self.resblocks = nn.Sequential(
        *[
            ResBlock(
                in_chans=conv_channels[-1],
                out_chans=conv_channels[-1],
                nonlinearity=nonlinearity,
                dropout=dropout_p
            ) for _ in range(n_blocks)
        ]
    )

    #output of resblocks is 16 16x16 images

    # Add final layers
    self.h2 = nn.Sequential()
    self.h2.add_module(
        name='final_batch_norm',
        module=nn.BatchNorm2d(
            num_features=conv_channels[-1]
        )
    )
    self.h2.add_module(
        name='final_relu',
        module=nn.ReLU()
    )
    self.h2.add_module(
        name = 'dropout_head',
        module=nn.Dropout(dropout_h)
    )
    self.h2.add_module(
        name = 'gap',
        module=nn.AvgPool2d(2)
    )

```

```

    )
    self.h2.add_module(
        name = 'flatten',
        module=nn.Flatten()
    )

    #output is 16 8x8 images
    # 16 comes from conv_channels[-1]
    # 8x8 comes from input_dim // 4
    fc_in = conv_channels[-1] * (input_dim//4)**2

    for i in range(len(fc_channels)):

        self.h2.add_module(
            name=f'fc{i}',
            module=nn.Linear(
                in_features=fc_in if i == 0 else fc_channels[i-1],
                out_features=fc_channels[i]
            )
        )
        if i < len(fc_channels) - 1:
            self.h2.add_module(
                name = f'fc_activation{i}',
                module=nn.ReLU()
            )

    self.h2.add_module('softmax', nn.Softmax(dim=1))
def forward(self, x):
    #print(f"h1_in: {x.shape}")
    out = self.h1(x)
    #print(f"res_block_in: {out.shape}")
    out = self.resblocks(out)
    #print(f"h2_in: {out.shape}")
    out = self.h2(out)
    return out

```

```

[5]: from torch.profiler import profile, record_function, ProfilerActivity
try:
    del train_loader
    del test_loader
    del model_1a
    del model_1b
    del resnet
    del train_loader_cuda
    del test_loader_cuda
except:
    pass

```

```

# Reset CUDA context
start = time.time()
torch.cuda.empty_cache()
torch.cuda.reset_peak_memory_stats()

gc.collect()

cifar10_train = datasets.CIFAR10(data_path, train=True, download=dl,
    ↪transform=transform)
cifar10_test = datasets.CIFAR10(data_path, train=False, download=dl,
    ↪transform=transform)

batch_size = int(2**11)
workers = 12
cpu_prefetch = 39
gpu_prefetch = 28

print('begin init train_loader')
train_loader = DataLoader(
    cifar10_train,
    batch_size=batch_size,
    shuffle=True,
    num_workers=workers,
    prefetch_factor=cpu_prefetch,
    pin_memory=True
)

X_batch = next(iter(train_loader))[0]
dtype_size = X_batch.element_size()
print(f"Batch Size: {X_batch.element_size() * X_batch.nelement() / 1024**2}
    ↪MiB")

print('begin init fetcher')
train_loader_cuda = CudaDataPrefetcher(
    data_iterable = train_loader,
    device = torch.device('cuda'),
    num_prefetch_batches=gpu_prefetch
)

test_loader = DataLoader(cifar10_test, batch_size=len(cifar10_test),
    ↪shuffle=True, num_workers=workers, pin_memory=True, prefetch_factor=1)
test_loader_cuda = CudaDataPrefetcher(
    data_iterable = test_loader,
    device = torch.device('cuda'),
    num_prefetch_batches=1
)

```

```

resnet = ResNet(
    input_dim = 32,
    conv_channels=[16,16],
    n_blocks = 10,
    fc_channels=[16,10],
    dropout_h = 0.6,
    dropout_p = 0.4
).to(device=device)
print(f"Init time: {(time.time() - start):.2f} seconds")
params = sum(p.numel() for p in resnet.parameters() if p.requires_grad)
print(f"Total parameters: {params}")
resnet.train_model(
    epochs=200,
    train_loader=train_loader_cuda,
    train_len=len(train_loader),
    test_loader=test_loader_cuda,
    test_len=len(test_loader),
    test_size=len(cifar10_test),
    loss_fn=nn.CrossEntropyLoss(),
    optimizer = torch.optim.Adam,
    optimizer_kwargs={'lr': 2e-3, 'weight_decay': 7e-3},
    print_epoch = 1,
    sched_factor = 0.6,
    sched_patience = 15
)

resnet.plot_training("ResNet Training Curves")

```

```

begin init train_loader
Batch Size: 24.0 MiB
begin init fetcher
Init time: 4.00 seconds
Total parameters: 66090
Training ResNet

```


	Epoch		Epoch Time (s)		Training Loss		Test Loss
	Overfit (%)		Accuracy (%)		Δ Accuracy (%)		GPU Memory
(GiB)							

	0		7.758229		2.262175		2.303453
	1.824688		9.999999		0.000000		4.617005

	1	5.941853	2.131489	2.312773
	8.505051	10.349999	3.500000	5.320131
	2	5.965214	2.039086	2.239328
	9.820173	18.820000	81.835754	6.021304
	3	5.957030	1.999411	2.137333
	6.898159	32.960001	75.132843	6.722476
	4	5.937284	1.963828	2.014531
	2.581829	47.529998	44.205090	7.423647
	5	5.978225	1.935574	2.026971
	4.721979	46.619999	-1.914578	8.124819
	6	5.998431	1.916032	2.005815
	4.685873	46.309999	-2.566800	8.825991
	7	5.885165	1.904790	1.935971
	1.637028	55.320001	16.389654	9.527163
	8	5.937993	1.893218	1.929178
	1.899396	55.729997	0.741136	10.228335

	9	5.912222	1.880680	1.889912
	0.490889	59.649998	7.033916	10.929507
	10	5.891649	1.870847	1.926282
	2.963132	54.809999	-8.113996	11.630679
	11	5.917224	1.862037	1.985746
	6.643737	47.759998	-19.932943	12.331851
	12	5.954925	1.852967	1.968036
	6.210001	50.389999	-15.523887	13.033022
	13	5.949610	1.850348	1.894329
	2.376869	58.849996	-1.341159	13.734194
	14	5.927939	1.838039	1.884981
	2.553910	58.950001	-1.173507	14.435366
	Epoch	Epoch Time (s)	Training Loss	Test Loss
	Overfit (%)	Accuracy (%)	Δ Accuracy (%)	GPU Memory
	(GiB)			
	15	5.978298	1.835852	1.916794
	4.408997	55.689996	-6.638729	15.136538

	16	5.931443	1.831612	1.992336
	8.774986	46.829998	-21.492037	15.136538
	17	5.936951	1.830642	1.878986
	2.640804	60.289997	1.072924	15.136538
	18	6.016370	1.825897	1.834697
	0.481930	63.940001	6.054079	15.136538
	19	5.955230	1.824874	1.879447
	2.990498	58.520001	-8.476695	15.136538
	20	5.945711	1.819507	1.890054
	3.877305	58.429998	-8.617457	15.136538
	21	5.960409	1.816705	1.909476
	5.106515	56.019998	-12.386617	15.136538
	22	5.979871	1.816134	1.913466
	5.359302	54.909998	-14.122619	15.136538
	23	5.943759	1.814242	1.888312
	4.082688	58.359998	-8.726934	15.136538

	24	5.989490	1.812754	1.932691
	6.616277	53.270000	-16.687519	15.136538
	25	5.976249	1.811979	1.853341
	2.282714	61.919999	-3.159214	15.136538
	26	5.926868	1.805859	1.890505
	4.687298	57.969999	-9.336880	15.136538
	27	5.957161	1.808381	1.828712
	1.124280	64.550000	0.954019	15.136538
	28	5.978532	1.808984	1.878929
	3.866507	58.649999	-9.140203	15.136538
	29	5.943629	1.806721	1.885318
	4.350254	58.569998	-9.264141	15.136538
	Epoch	Epoch Time (s)	Training Loss	Test Loss
	Overfit (%)	Accuracy (%)	Δ Accuracy (%)	GPU Memory
	(GiB)			
	30	5.946413	1.802864	1.968687
	9.197787	48.940000	-24.182804	15.136538

	31	5.973595	1.804069	1.820951
	0.935777	65.279996	1.130900	15.136538
	32	5.941212	1.800107	1.835910
	1.988929	63.870001	-2.159920	15.136538
	33	5.957007	1.798617	1.939540
	7.835098	52.329999	-19.837620	15.136538
	34	5.956561	1.798999	1.824884
	1.438882	64.569998	-1.087620	15.136538
	35	6.000816	1.799080	1.926372
	7.075355	53.549999	-17.968748	15.136538
	36	5.929878	1.797763	1.850310
	2.922882	61.909997	-5.162376	15.136538
	37	5.952627	1.798736	1.868542
	3.880836	60.339999	-7.567398	15.136538
	38	5.983609	1.794371	1.869053
	4.161981	59.729999	-8.501835	15.136538

	39	5.951901	1.791374	1.866068
	4.169658	60.249996	-7.705270	15.136538
	40	6.068545	1.793264	1.898383
	5.861915	56.849998	-12.913601	15.136538
	41	6.035099	1.790199	1.822050
	1.779197	64.849997	-0.658701	15.136538
	42	5.953965	1.793583	1.866762
	4.080087	60.039997	-8.026960	15.136538
	43	6.000158	1.793789	1.888961
	5.305665	57.660002	-11.672788	15.136538
	44	6.012037	1.788270	1.844360
	3.136523	62.459999	-4.319849	15.136538
	Epoch	Epoch Time (s)	Training Loss	Test Loss
	Overfit (%)	Accuracy (%)	Δ Accuracy (%)	GPU Memory
	(GiB)			
	45	6.064523	1.786979	1.848547
	3.445322	61.769998	-5.376836	15.136538

	46		6.086365		1.789465		1.920969
	7.348784		54.629999		-16.314335		15.136538
							</

	54	5.964084	1.764875	1.791243
	1.494045	67.729998	-0.235679	15.136538
	55	5.955197	1.766104	1.816597
	2.859046	65.169996	-4.006487	15.136538
	56	6.029804	1.763746	1.872248
	6.151809	58.849996	-13.315663	15.136538
	57	5.953676	1.762739	1.792457
	1.685915	67.839998	-0.073652	15.136538
	58	5.990416	1.761238	1.820805
	3.382094	64.569998	-4.890267	15.136538
	59	5.994577	1.761105	1.813354
	2.966815	65.230000	-3.918103	15.136538
	Epoch	Epoch Time (s)	Training Loss	Test Loss
	Overfit (%)	Accuracy (%)	Δ Accuracy (%)	GPU Memory
	(GiB)			
	60	6.018641	1.758287	1.878296
	6.825368	59.749997	-11.989988	15.136538

	69	6.017951	1.723430	1.808294
	4.924157	66.399997	-6.898486	15.136538
	70	6.035763	1.725271	1.767110
	2.425069	70.190001	-1.584404	15.136538
	71	5.997301	1.720044	1.785120
	3.783398	68.059999	-4.570946	15.136538
	72	6.033640	1.718865	1.798987
	4.661331	66.780001	-6.365671	15.136538
	73	6.025136	1.722847	1.744101
	1.233708	72.560000	1.738647	15.136538
	74	6.044716	1.716789	1.828936
	6.532368	63.679999	-12.238149	15.136538
	Epoch	Epoch Time (s)	Training Loss	Test Loss
	Overfit (%)	Accuracy (%)	Δ Accuracy (%)	GPU Memory
	(GiB)			
	75	6.046995	1.724921	1.802198
	4.480049	66.319996	-8.599786	15.136538

	76		6.073751		1.722809		1.788122
	3.791047		67.979997		-6.312021		15.136538
	77		5.978110		1.715931		1.788931
	4.254255		68.110001		-6.132854		15.136538
	78		6.042846		1.714050		1.765153
	2.981410		70.389998		-2.990632		15.136538
	79		5.989166		1.709710		1.810159
	5.875241		65.799999		-9.316429		15.136538
	80		6.026500		1.713593		1.827208
	6.630245		63.989997		-11.810920		15.136538
	81		6.064908		1.713212		1.820815
	6.280786		64.289999		-11.397466		15.136538
	82		6.053454		1.713486		1.755603
	2.457978		71.099997		-2.012133		15.136538
	83		6.030574		1.709514		1.789637
	4.686906		67.890000		-6.436053		15.136538

	84	6.010733	1.713726	1.790577
	4.484475	67.530000	-6.932194	15.136538
	85	5.997445	1.712908	1.741797
	1.686579	72.700000	0.192943	15.136538
	86	6.078503	1.707408	1.794512
	5.101524	67.109996	-7.689139	15.136538
	87	6.048886	1.713436	1.799901
	5.046287	66.710001	-8.239338	15.136538
	88	6.026438	1.711424	1.845643
	7.842492	61.939996	-14.800554	15.136538
	89	6.002117	1.707037	1.757974
	2.983910	71.410000	-1.774415	15.136538
	Epoch	Epoch Time (s)	Training Loss	Test Loss
	Overfit (%)	Accuracy (%)	Δ Accuracy (%)	GPU Memory
	(GiB)			
	90	5.973038	1.710023	1.775080
	3.804483	69.150001	-4.883080	15.136538

	91	6.030484	1.706634	1.743197
	2.142406	72.649997	-0.068779	15.136538
	92	6.044240	1.706756	1.857970
	8.859737	60.549998	-16.712519	15.136538
	93	6.025276	1.709503	1.789289
	4.667165	67.949998	-6.533702	15.136538
	94	6.065005	1.709287	1.730637
	1.249039	73.869997	1.609350	15.136538
	95	6.063215	1.706813	1.837258
	7.642605	62.809998	-14.972248	15.136538
	96	6.046596	1.709644	1.749821
	2.349999	72.299999	-2.125353	15.136538
	97	6.041147	1.707956	1.765295
	3.357140	70.229995	-4.927578	15.136538
	98	6.140745	1.705159	1.762154
	3.342478	70.159996	-5.022339	15.136538

	99	5.942571	1.704508	1.816739
	6.584328	64.919996	-12.115881	15.136538
	100	5.976495	1.704234	1.753127
	2.868941	71.859998	-2.720996	15.136538
	101	6.089538	1.705383	1.850588
	8.514546	61.739999	-16.420738	15.136538
	102	6.001499	1.704032	1.744620
	2.381902	72.499996	-1.854611	15.136538
	103	6.000335	1.703446	1.796068
	5.437298	67.570001	-8.528491	15.136538
	104	5.989391	1.702863	1.846086
	8.410723	61.839998	-16.285366	15.136538
	Epoch	Epoch Time (s)	Training Loss	Test Loss
	Overfit (%)	Accuracy (%)	Δ Accuracy (%)	GPU Memory
	(GiB)			
	105	6.053767	1.704828	1.784842
	4.693377	68.210000	-7.662106	15.136538

	106	5.995612	1.704311	1.820918
	6.841868	63.989997	-13.374848	15.136538
	107	6.119347	1.703072	1.791096
	5.168544	67.799997	-8.217138	15.136538
	108	5.988535	1.702467	1.789317
	5.101415	68.229997	-7.635035	15.136538
	109	6.011374	1.704103	1.806670
	6.018831	65.700001	-11.059965	15.136538
	110	6.016177	1.705480	1.752431
	2.752946	71.799999	-2.802217	15.136538
	111	6.034240	1.686944	1.737605
	3.003156	73.249996	-0.839314	15.136538
	112	6.049294	1.681191	1.718573
	2.223586	74.759996	1.204818	15.136538
	113	6.035938	1.682705	1.719535
	2.188763	75.009996	0.334403	15.136538

	114	6.005780	1.682150	1.720372
	2.272208	74.759996	-0.333289	15.136538
	115	6.040393	1.683993	1.763142
	4.700065	70.400000	-6.145842	15.136538
	116	6.026464	1.679085	1.732503
	3.181357	73.899996	-1.479802	15.136538
	117	6.116028	1.679812	1.728820
	2.917485	73.769999	-1.653109	15.136538
	118	6.095558	1.680922	1.724223
	2.576033	74.169999	-1.119846	15.136538
	119	6.083001	1.678301	1.736588
	3.472963	73.119998	-2.519661	15.136538
	Epoch	Epoch Time (s)	Training Loss	Test Loss
	Overfit (%)	Accuracy (%)	Δ Accuracy (%)	GPU Memory
	(GiB)			
	120	6.071417	1.680773	1.722848
	2.503350	74.419999	-0.786557	15.136538

	121	6.101312	1.682564	1.713852
	1.859566	75.769997	1.013200	15.136538
	122	6.108806	1.682129	1.731231
	2.919064	73.869997	-2.507588	15.136538
	123	6.083366	1.681966	1.703076
	1.255068	76.769996	1.319782	15.136538
	124	6.085177	1.679816	1.698992
	1.141518	76.580000	-0.247487	15.136538
	125	6.054791	1.680250	1.748276
	4.048547	71.819997	-6.447830	15.136538
	126	6.054725	1.681701	1.729124
	2.819957	74.149996	-3.412791	15.136538
	127	6.105957	1.682942	1.785550
	6.096926	67.839998	-11.632145	15.136538
	128	6.027437	1.680754	1.787973
	6.379204	67.739999	-11.762403	15.136538

	129	6.026318	1.681985	1.742985
	3.626660	72.439998	-5.640221	15.136538
	130	6.067482	1.681037	1.717516
	2.170037	74.799997	-2.566106	15.136538
	131	6.053882	1.680817	1.752610
	4.271327	71.630001	-6.695317	15.136538
	132	6.073949	1.681933	1.761947
	4.757300	70.599997	-8.036992	15.136538
	133	6.100069	1.679813	1.729257
	2.943421	73.749995	-3.933829	15.136538
	134	6.060847	1.681197	1.747969
	3.971707	72.179997	-5.978897	15.136538
	Epoch	Epoch Time (s)	Training Loss	Test Loss
	Overfit (%)	Accuracy (%)	Δ Accuracy (%)	GPU Memory
	(GiB)			
	135	6.092427	1.676345	1.728331
	3.101143	73.869997	-3.777515	15.136538

	136		6.111445		1.677436		1.710195
	1.952925		75.689995		-1.406800		15.136538
	137		6.154238		1.680783		1.780700
	5.944660		68.439996		-10.850594		15.136538
	138		6.104064		1.680163		1.767590
	5.203455		69.580001		-9.365631		15.136538
	139		6.064308		1.679928		1.724739
	2.667443		74.210000		-3.334630		15.136538
	140		6.095605		1.670509		1.695053
	1.469260		77.579999		1.055104		15.136538
	141		6.075877		1.665254		1.685887
	1.239022		78.099996		0.670272		15.136538
	142		6.111793		1.662272		1.704352
	2.531463		76.239997		-2.381561		15.136538
	143		6.030466		1.663150		1.701710
	2.318458		76.569998		-1.959024		15.136538

	144	6.110031	1.664230	1.696805
	1.957377	76.959997	-1.459665	15.136538
	145	6.130999	1.662143	1.703371
	2.480386	76.400000	-2.176692	15.136538
	146	6.076478	1.659864	1.729352
	4.186393	73.600000	-5.761838	15.136538
	147	6.191991	1.659926	1.678746
	1.133810	79.029995	1.190780	15.136538
	148	6.065806	1.664576	1.687991
	1.406648	78.039998	-1.252686	15.136538
	149	6.120574	1.660582	1.703497
	2.584342	76.519996	-3.176008	15.136538
	Epoch	Epoch Time (s)	Training Loss	Test Loss
	Overfit (%)	Accuracy (%)	Δ Accuracy (%)	GPU Memory
	(GiB)			
	150	6.083641	1.660676	1.691585
	1.861263	77.609998	-1.796783	15.136538

	151	6.159505	1.661676	1.726976
	3.929745	74.030000	-6.326706	15.136538
	152	6.116043	1.662686	1.677465
	0.888861	78.789997	-0.303680	15.136538
	153	6.153685	1.661415	1.669934
	0.512795	79.769999	0.936358	15.136538
	154	6.083883	1.660522	1.722128
	3.710041	74.159998	-7.032721	15.136538
	155	6.110673	1.661754	1.710175
	2.913857	75.610000	-5.214992	15.136538
	156	6.137404	1.660498	1.689023
	1.717822	77.899998	-2.344240	15.136538
	157	6.167748	1.659847	1.681879
	1.327353	78.389996	-1.729977	15.136538
	158	6.080560	1.659312	1.732317
	4.399757	73.509997	-7.847564	15.136538

	159	6.110744	1.659585	1.717616
	3.496686	74.640000	-6.430988	15.136538
	160	6.144540	1.661628	1.708873
	2.843304	75.439996	-5.428110	15.136538
	161	6.116296	1.661444	1.721828
	3.634423	74.469995	-6.644106	15.136538
	162	6.133440	1.659592	1.694767
	2.119488	76.919997	-3.572774	15.136538
	163	6.133320	1.661856	1.685046
	1.395393	78.200001	-1.968156	15.136538
	164	6.137462	1.661753	1.698507
	2.211732	76.719999	-3.823492	15.136538
	Epoch	Epoch Time (s)	Training Loss	Test Loss
	Overfit (%)	Accuracy (%)	Δ Accuracy (%)	GPU Memory
	(GiB)			
	165	6.100959	1.660505	1.710490
	3.010189	75.489998	-5.365427	15.136538

	166	6.162635	1.660260	1.731846
	4.311771	73.479998	-7.885171	15.136538
	167	6.242150	1.659401	1.695021
	2.146534	77.279997	-3.121477	15.136538
	168	6.126612	1.662474	1.725369
	3.783221	73.850000	-7.421335	15.136538
	169	6.159417	1.660429	1.695292
	2.099681	77.079999	-3.372194	15.136538
	170	6.097894	1.653447	1.680014
	1.606785	78.639996	-1.416576	15.136538
	171	6.127400	1.648021	1.669070
	1.277208	79.850000	0.100290	15.136538
	172	6.093162	1.646899	1.670808
	1.451743	79.439998	-0.513466	15.136538
	173	6.118891	1.646340	1.681919
	2.161102	78.469998	-1.728244	15.136538

	174	6.095659	1.646992	1.669149
	1.345276	79.749995	-0.125241	15.136538
	175	6.131203	1.648076	1.677630
	1.793243	78.799999	-1.314967	15.136538
	176	6.111234	1.649875	1.681801
	1.935102	78.560001	-1.615528	15.136538
	177	6.161559	1.647128	1.718898
	4.357234	74.390000	-6.837821	15.136538
	178	6.125992	1.647206	1.660054
	0.780025	80.739999	1.114588	15.136538
	179	6.156801	1.645703	1.675382
	1.803398	79.189998	-1.919744	15.136538
	Epoch	Epoch Time (s)	Training Loss	Test Loss
	Overfit (%)	Accuracy (%)	Δ Accuracy (%)	GPU Memory
	(GiB)			
	180	6.101541	1.646676	1.699421
	3.203148	76.969999	-4.669309	15.136538

	181		6.070375		1.646551		1.670855
	1.476050		79.560000		-1.461480		15.136538
	182		6.133426		1.645668		1.692414
	2.840550		77.370000		-4.173890		15.136538
	183		6.103067		1.646507		1.673047
	1.611929		79.369998		-1.696806		15.136538
	184		6.220975		1.645006		1.687924
	2.609001		77.609998		-3.876642		15.136538
	185		6.173976		1.645559		1.670266
	1.501455		79.539996		-1.486255		15.136538
	186		6.171250		1.644896		1.676799
	1.939509		78.759998		-2.452317		15.136538
	187		6.168568		1.644711		1.667325
	1.374997		79.879999		-1.065147		15.136538
	188		6.133537		1.644299		1.690048
	2.782248		77.759999		-3.690859		15.136538

	189	6.135493	1.645682	1.683321
	2.287131	78.410000	-2.885805	15.136538
	190	6.136440	1.646676	1.672363
	1.559951	79.189998	-1.919744	15.136538
	191	6.104340	1.644953	1.697103
	3.170348	76.800001	-4.879859	15.136538
	192	6.190171	1.642390	1.673998
	1.924519	79.119998	-2.006441	15.136538
	193	6.137682	1.644042	1.684601
	2.467033	78.109998	-3.257371	15.136538
	194	6.132952	1.644127	1.669210
	1.525608	79.549998	-1.473868	15.136538
	Epoch	Epoch Time (s)	Training Loss	Test Loss
	Overfit (%)	Accuracy (%)	Δ Accuracy (%)	GPU Memory
	(GiB)			
	195	6.213376	1.641013	1.665620
	1.499475	80.070001	-0.829822	15.136538

	196		6.159393		1.635479		1.667431
	1.953695		80.039996		-0.866984		15.136538

	197		6.220014		1.636976		1.664628
	1.689192		80.129999		-0.755511		15.136538

	198		6.141409		1.638255		1.662055
	1.452788		80.409998		-0.408721		15.136538

	199		6.201807		1.636783		1.666127
	1.792761		79.899997		-1.040379		15.136538

Training Time: 667.1745541095734 seconds

